

This thesis was submitted to the Chair of Machine Learning and Reasoning.

Semantic Partitioning for Partially Observable Deterministic Task and Motion Planning

Master Thesis

Presented by

Hani Alassiri Alhabboub
422433

Supervised by Prof. Hector Geffner, Ph.D.

1st Examiner Prof. Hector Geffner, Ph.D.

2nd Examiner Prof. Christopher Morris

Aachen, June 10, 2026

Everything is Computer

DJT

Abstract

This thesis extends classical Task and Motion Planning (TAMP) to partial observability by integrating a Semantic Boxel representation—task-relevant cuboids that concentrate resolution on the objects and the regions they occlude, covering the free space with a few coarse cells—into a PDDLStream-based planner. The robot’s belief about which Boxels are occupied is encoded as Know-If literals—one per Boxel, marking whether its contents are known—so a classical planner can reason about hidden objects and plan sensing actions in the same search, without enumerating an exhaustive state space. On three tabletop tasks evaluated against a uniform-grid baseline at the same cell scale, the adaptive partition produces an order of magnitude fewer cells, cuts per-call planning time accordingly, and solves find-and-tray-stack scenes that the uniform baseline effectively cannot (39.8 % vs 1.3 % success).

Contents

1	Introduction	1
1.1	Contributions of This Thesis	3
1.2	Thesis Outline	3
2	Background	4
2.1	AI Planning Fundamentals	4
2.1.1	State-Space Models in Classical Planning	4
2.2	Planning under Partial Observability	6
2.2.1	Belief States and K-Literals	6
2.2.2	Probabilistic vs. Deterministic Partial Observability	7
2.3	Task and Motion Planning (TAMP)	7
2.3.1	PDDLStream for TAMP	7
2.4	Spatial Representation for Robotic Planning	9
2.4.1	Discretization of Continuous Space	9
3	Related Work	11
3.1	TAMP under Partial Observability	11
3.1.1	POMDP-based TAMP Solutions	11
3.1.2	Partially Grounded Plans in TAMP	11
3.1.3	Modern Approaches to TAMP under Uncertainty	12
3.1.4	Belief-Space Planning and Replanning	13
3.1.5	Partially Observable Deterministic Planning	13
3.1.6	Summary and Research Gap	14
3.2	Spatial Belief Representation in TAMP	14
3.2.1	Limitations of Geometric Discretization for POD-TAMP	14
3.2.2	Limitations of Static Semantic Abstractions	15
4	Approach	16
4.1	Overview	16
4.2	Adaptive Semantic Discretization: Generating Boxels	16
4.3	Belief over Boxels with Know-If Fluents	19
4.4	Semantic POD-TAMP Model Definition	20
4.5	PDDLStream Integration	21
4.5.1	PDDL Fluents for Belief Representation	21

4.5.2	The Sensing Action	21
5	Evaluation	24
5.1	Experimental Setup	24
5.2	Evaluation Metrics	26
5.3	Baselines for Comparison	27
5.4	Results	28
5.4.1	Anytime Solve Rate	28
5.4.2	Success Rate and Scene Difficulty	29
5.4.3	Planning Time and Scene Difficulty	31
5.4.4	Representation Compactness	31
5.4.5	Free-Space Resolution	32
5.4.6	Comparison with TAMPURA	33
5.4.7	Failure Modes	34
6	Discussion	36
6.1	Adaptive vs Uniform Partitioning	36
6.2	The 5 cm Leaf-Floor Variant	36
6.3	Architectural Comparison with TAMPURA	37
6.4	Threats to Validity	38
6.5	Limitations and Accepted Simplifications	39
7	Conclusion	42
7.1	Future Work	43
A	PDDL Specification	44
A.1	Domain	44
A.2	Streams	49
	List of Acronyms	51
	List of Symbols	52
	List of Figures	54
	List of Tables	57
	List of Listings	58
	List of References	59

1 Introduction

We want robots that can carry out everyday tasks on their own—fetching an object, clearing a table—tasks that demand both multi-step reasoning and careful physical manipulation. Task and motion planning (TAMP) studies how to do exactly this [7, 10]. TAMP connects a high-level goal (e.g., “fetch the blue cup from the kitchen”) to the actual motions the robot must perform to achieve it. TAMP combines discrete task planning with continuous motion planning: symbolic planners generate high-level sequences of abstract actions—what is known as a task plan. However, symbolic planning alone is insufficient for real-world robotic execution. Because the symbolic plan ignores geometry, it can call for actions the robot cannot physically carry out. Specifically, we must determine whether the proposed symbolic actions are feasible within the robot’s continuous geometric and kinematic constraints, and if so, identify the concrete parameters required to execute them—such as object poses, grasp configurations, and collision-free trajectories. For instance, a symbolic instruction like (*pick A*) is only viable if object A is actually reachable, graspable, and a feasible motion exists to carry out the action.

One widely used framework for this is PDDLStream [10]. PDDLStream extends the classical Planning Domain Definition Language (PDDL) [23] by incorporating streams—lazily-evaluated, declarative procedures that can interface with external solvers. These streams are capable of generating or validating continuous parameters on demand, allowing the symbolic planner to retrieve geometric or kinematic information only as needed. This simplifies the coupling between high-level symbolic reasoning (what to do) and low-level geometric feasibility (how to do it), making the planning process both flexible and efficient.

However, the original PDDLStream framework assumes full observability and deterministic action outcomes—assumptions that rarely hold in practice. In practice, robots frequently operate under partial observability, due to occlusions, limited sensor coverage, or dynamically changing environments. In such settings, crucial information—such as the location of a target object or the presence of an obstacle—may be unknown at the start, so the robot has to actively look for it while carrying out the task (Figure 1.1).




Figure 1.1: A 7-DOF Franka Panda at a cluttered tabletop. The target—the cyan cube behind the green cube—is hidden from the overhead camera and must be located before it can be retrieved. The corner panels are the fixed overhead camera’s synthetic RGB (top) and depth (bottom) views.

Addressing this limitation is the central focus of this thesis: extending PDDLStream to operate under partial observability. This includes the ability to represent and reason over belief states that capture uncertainty in continuous spatial variables. Naive representations—such as a uniform voxel grid over 3D space—make the belief space blow up in size and quickly become impractical to plan with. A useful belief representation must therefore be detailed enough to be expressive, yet small enough to plan with quickly. The core research question is: *How can we formally define and integrate a belief representation based on semantic partitioning into a POD-TAMP framework?* The stance of this thesis is simple: *space is part of the planning problem.* Where a hidden object might be, and what blocks the view of it, are not details to be collapsed into a probabilistic occupancy grid but structure the planner should weigh directly. Here, we present a novel approach to discretizing the workspace around the objects the robot detects, the regions they occlude, and the free space between them. The unit of this discretization is the Semantic Boxel: a task-relevant, occlusion-aware cuboid, built from objects’ known 3D models, that the robot uses to represent its belief. Unlike a dense, uniform discretization, this provides a structured, scalable way to represent and reason about spatial uncertainty.

This thesis develops a system in which a robot, facing uncertainty about object locations, can plan and execute a sequence of information-gathering and manipulation actions. It reasons over abstract beliefs about which Boxels objects might occupy, using PDDLStream to connect this reasoning to continuous motion and manipulation. Sensing is a symbolic action evaluated against a fixed overhead camera, not a viewpoint the robot moves to; and because the underlying planner is classical, partial observability is handled by optimistic sensing with reactive replanning rather than contingent planning—a sensing action is assumed to succeed, and the system replans if execution shows otherwise. The result is a framework that makes

planning under spatial uncertainty simpler and more reliable. The approach is demonstrated on tabletop hidden-object retrieval and stacking; the partial-observability mechanism is the contribution, and longer task horizons are a natural extension.

1.1 Contributions of This Thesis

The central contribution is a *representation*: each volume the robot cannot see becomes a named region that a classical, partially observable deterministic (POD) planner reasons over and resolves by sensing. The three components below deliver this representation; the evaluation correspondingly isolates it by ablation (the same planner on the adaptive partition against a uniform grid), rather than claiming a benchmark win, and is conducted under oracle perception (Section 6.5).

- The *Semantic Boxel* abstraction: an occlusion-aware, object-centric partitioning of the workspace that concentrates resolution on the task-relevant regions—the objects themselves and the volumes they occlude—while covering the remaining free space with a few coarse cells, rather than partitioning all space uniformly.
- A Semantic POD-TAMP model that uses Know-If Fluents over Boxels as a compact belief representation, letting a classical PDDL planner reason about partial observability via optimistic sensing with reactive replanning.
- A PDDLStream realization of this model that combines the discrete Boxel belief with stream-certified continuous geometry, so that information-gathering and manipulation actions interleave within a single planning domain.

1.2 Thesis Outline

Chapter 2 provides essential background on classical planning, reasoning under uncertainty, and the PDDLStream framework that this work builds upon. Chapter 3 reviews related work in TAMP under partial observability, belief representation, and adjacent planning paradigms. Chapter 4 details the methodology, including the formalization of Boxels through adaptive semantic discretization, the Semantic POD-TAMP model, and the PDDLStream integration. Chapter 5 presents the experimental setup, the metrics and baselines, and the headline evaluation results. Chapter 6 interprets those results and consolidates the accepted simplifications. Chapter 7 summarizes the contributions and outlines directions for future work.

2 Background

This chapter builds up the pieces the thesis depends on, ending with PDDLStream and how 3D space is represented for planning.

2.1 AI Planning Fundamentals

Artificial intelligence (AI) planning produces a sequence of actions—a plan—that gets from a given initial state to a specified goal [13].

2.1.1 State-Space Models in Classical Planning

The formal basis for classical planning problems is a *state-space model* [12, 13, 17]. A state-space model defines every possible configuration of the world and how actions move between configurations.

More formally, a deterministic state-space model can be described as a tuple $M = \langle S, s_0, S_G, Act, A, f \rangle$, where:

- S : A set representing all possible *states* the world can be in. For example, in a simple blocksworld problem with two blocks A and B and a table, a state could be "Block A is on Block B, and Block B is on the table, and the robot hand is empty."
- $s_0 \in S$: The specific *initial state* from which planning begins. E.g., "Block A is on the table, Block B is on the table, hand empty."
- $S_G \subseteq S$: The set of *goal states*. Any state in this set satisfies the desired objective. E.g., any state where "Block A is on Block B."
- Act : A set of all possible *actions* the agent can perform. E.g., *pick blockX*, *stack blockX on blockY*.
- $A(s) \subseteq Act$: A function that, for any given state $s \in S$, returns the subset of actions from Act that are *applicable* (can be executed) in state s . For example, *pick blockA* is applicable only if Block A is clear (nothing on top of it) and the robot hand is empty.
- $f(a, s) \rightarrow s'$: A deterministic *state transition function*. Given a state s and an applicable action $a \in A(s)$, $f(a, s)$ returns the unique successor state s' that results from executing action a in state s . For example, if s is "A on table, B on table, hand empty" and a is *pick blockA*, then $s' = f(a, s)$ would be "B on table, hand holding A."

A *solution* or *plan* is then a sequence of actions $\sigma = [a_0, a_1, \dots, a_n]$ such that if s_i is the state before action a_i , then $a_i \in A(s_i)$, $s_{i+1} = f(a_i, s_i)$, and the final state s_{n+1} is in S_G . Often, we

are interested in an *optimal plan*, which could be the shortest sequence of actions or one that minimizes some associated cost.

Compact representation: STRIPS, and PDDL. Explicitly enumerating all possible states S is infeasible for all but the simplest problems. Therefore, planning problems are typically described in a more compact, *factored* representation using a planning language. In this representation, a state is defined by a set of propositions (or ground atoms) that are currently true.

A foundational language for this is **Stanford Research Institute Problem Solver (STRIPS)** [8]. In STRIPS:

- **States** are represented as sets of true propositions (e.g., (on A B), (ontable B), (handempty)).
- **Actions** are defined by:
 - *Preconditions*: A set of propositions that must be true in the current state for the action to be applicable.
 - *Effects*: Specified as two lists:
 - * *Add-list*: A set of propositions that become true after the action is executed.
 - * *Delete-list*: A set of propositions that become false (are no longer true) after the action.
- The transition function $f(a, s)$ is implicitly defined: the new state s' is obtained by taking state s , removing all propositions in the delete-list of a , and then adding all propositions in the add-list of a .
- The initial state (s_0) is a set of initially true propositions, and the goal (S_G) is specified by a set of propositions that must be true in any goal state.

The **PDDL** [23] builds upon and extends the STRIPS representation, providing a more expressive and standardized syntax for a wider range of classical planning problems. In PDDL:

- **Domain File** defines:
 - *Types*: Categories for objects (e.g., block, robot).
 - *Predicates*: Relations or properties (e.g., (on ?x - block ?y - block)).
 - *Action Schemas*: Templates for actions, generalizing STRIPS operators. For instance:

```

1          (:action stack
2              :parameters (?obj1 - block ?obj2 - block)
3              :precondition (and (holding ?obj1) (clear ?obj2))
4              :effect (and (on ?obj1 ?obj2) (not (holding ?obj1))
5                          (not (clear ?obj2)) (handempty) (clear
6                              ?obj1))
9          )

```

In STRIPS terms, this action's delete-list contains (holding ?obj1) and (clear ?obj2), while its add-list contains (on ?obj1 ?obj2), (handempty), and (clear ?obj1).

- **Problem File** specifies:
 - *Objects*: Specific objects (e.g., A, B, C - block).
 - *Initial State (Init)*: Ground atoms true at the start.
 - *Goal (G)*: Ground atoms that must be true in a goal state.

A PDDL problem $P = \langle \text{Domain}, \text{Instance} \rangle$ thus compactly encodes the state model $S(P)$. Ground actions (action schemas with parameters replaced by specific objects) form the set *Act*. More advanced PDDL versions also support features beyond STRIPS, such as conditional effects ((when <condition> <effect>)), numeric fluents, and axioms (derived predicates).

2.2 Planning under Partial Observability

Classical planning's assumption of full observability is often violated in robotics. When an agent has incomplete information about the true state of the world, it must reason about its *belief state*.

2.2.1 Belief States and K-Literals

A belief state represents the agent's current set of possible world states consistent with its history of actions and observations. We refer to the setting this thesis targets as *Partially Observable Deterministic (POD)* planning: the world itself behaves deterministically and the only uncertainty is the agent's incomplete knowledge of the initial situation, which sensing resolves. POD builds on [1, 5, 12]. The belief state b is explicitly a set of possible states. Actions transition this set, and observations prune it. A common way to reason about knowledge in POD planning is through K-literals [4, 12]. K-literals make the robot's knowledge about a proposition p explicit. For example:

- $K(p)$ signifies " p is known to be true."
- $K(\neg p)$ signifies " p is known to be false."

If neither $K(p)$ nor $K(\neg p)$ holds, the truth of p is unknown. Separately, p is "possibly true" whenever $K(\neg p)$ does not hold—which also covers the case where p is known to be true.

K-literals are the basis of a *translation* that compiles a partially observable problem into a classical one [12]. Each literal L of the original problem is replaced by the two fluents $K(L)$ and $K(\neg L)$; in the translated problem these are made true in the initial state only where the agent actually knows the corresponding value, and a literal whose value is unknown contributes neither. The translated problem therefore carries no uncertainty in its initial state and can be solved by an ordinary classical planner, which now searches for a plan that reaches a desired

state of knowledge—for example $K(\text{goal_achieved})$ —rather than a desired world state. Because the pair $K(L)$ and $K(\neg L)$ together encodes whether the agent knows the truth of L , we call such a knowledge fluent a *Know-If Fluent (KIF)*. KIFs are the belief representation used throughout this thesis (Chapter 4).

2.2.2 Probabilistic vs. Deterministic Partial Observability

POD planning differs from planning with partially observable Markov decision processes (POMDPs) [19]—a distinction that matters because we later compare our approach to POMDP-based TAMP. In POMDPs, the world is modeled as a Markov decision process (MDP) with probabilistic transitions and observations, and the belief state is a probability distribution over world states. While highly expressive, POMDPs are computationally very demanding. POD planning assumes outcomes and observations are deterministic but unknown. This makes it more tractable than POMDPs, and it fits robotic problems where the uncertainty comes from not knowing the world rather than from actions behaving randomly.

2.3 Task and Motion Planning (TAMP)

TAMP connects high-level task goals with a robot’s actual physical movements [9]. The core problem is ensuring that a plan is physically possible. High-level planners often ignore geometry, leading to plans that seem correct but can’t be executed. Yet, including all geometric details from the start makes the problem too complex to solve. Combining high-level logic with real-world geometry is the central difficulty of TAMP, and frameworks like PDDLStream exist to bridge it.

2.3.1 PDDLStream for TAMP

PDDLStream [10] is a TAMP framework built to bridge this symbolic-continuous gap. It extends the classical PDDL by introducing the core concept of **streams**.

Streams: declarative procedures for continuous parameters. A stream in PDDLStream is a procedure that generates or tests values for continuous parameters required by symbolic actions. These parameters might include object poses, robot configurations, grasp transformations, or motion trajectories. Streams have both a procedural and a declarative component:

- **Procedural Component (Conditional Generator):** This is the actual code (e.g., a Python function) that, given some input values (which can be discrete symbols or other continuous values), produces a (potentially infinite) sequence of output values. For example, an inverse kinematics (IK) stream, given an object, its target pose, and a grasp, would generate a sequence of robot configurations that achieve that end-effector pose. A collision-checking stream, given a trajectory, would test if it’s collision-free.

- **Declarative Component (Stream Specification):** This part is described in a PDDL-like syntax and informs the symbolic planner what the stream does, what inputs it requires, what outputs it produces, and, crucially, what logical conditions (fluents) are associated with these inputs and outputs. This specification includes:
 - **:inputs:** The typed parameters the stream procedure takes.
 - **:domain:** A set of PDDL facts (preconditions) that must be true for the stream to be applicable for a given set of inputs. For example, an IK stream might require that a valid pose and grasp for the object are already known.
 - **:outputs:** The typed parameters that the stream procedure generates.
 - **:certified:** A set of PDDL facts that are guaranteed to be true if the stream successfully produces a set of outputs for given inputs. These “certified fluents” are added to the planner’s state and represent new knowledge derived from the continuous domain. For instance, if an IK stream successfully finds a configuration ?q for picking object ?obj at pose ?p with grasp ?g, it might certify the fluent (KinSolution ?obj ?p ?g ?q).

PDDLStream employs a **lazy evaluation** strategy. Streams are not called exhaustively beforehand; instead, the symbolic planner calls them “on-demand” when it encounters an action whose preconditions require a continuous parameter to be instantiated or a continuous condition to be tested.

Example PDDLStream workflow. Consider a symbolic action (pick ?obj ?p ?g ?q ?t) (pick object ?obj from pose ?p using grasp ?g via configuration ?q along trajectory ?t). For this action to be applicable:

1. A stream might first be called to sample a stable ?p for ?obj on a surface. If successful, it certifies (AtPose ?obj ?p).
2. Then, another stream samples a grasp ?g for ?obj. If successful, it certifies (Grasp ?obj ?g).
3. An IK stream then takes ?obj, ?p, ?g as input and tries to find a robot configuration ?q. If successful, it certifies (KinSolution ?obj ?p ?g ?q).
4. Finally, a motion planning stream takes the current configuration and ?q as input to find a collision-free trajectory ?t, certifying (Motion ?q_start ?t ?q).

Only if all these streams succeed and certify their respective fluents will the preconditions of the symbolic (pick ...) action be met.

The overall PDDLStream algorithm often reduces a TAMP problem to a sequence of finite PDDL problems, iteratively calling streams to discover more certified fluents until a complete, grounded plan is found. The original PDDLStream framework, however, was primarily designed for fully observable and deterministic settings.

2.4 Spatial Representation for Robotic Planning

A robot needs some way to represent the 3D space around it and the objects in it. Representing space is harder under partial observability, because the robot must also represent what it does not yet know.

2.4.1 Discretization of Continuous Space

Robots operate in continuous space, but symbolic planners reason over discrete states — so the continuous space has to be discretized somehow.

Voxel grids. A common approach to represent 3D environments is to discretize them into a grid of uniform volumetric pixels, or **voxels**. Each voxel can store information, such as whether it is occupied, free, or unknown. While straightforward, uniform voxel grids can be memory-intensive and computationally expensive for large environments or high resolutions.

Octrees (Figure 2.1) [18] offer a more efficient hierarchical data structure for representing 3D space. An octree is a tree data structure in which each internal (non-leaf) node has exactly eight children, while leaf nodes have none. It starts with a single large voxel encompassing the entire space of interest. If this voxel is not homogeneous (e.g., it contains both free and occupied space), it is recursively subdivided into eight smaller child voxels (octants).




Figure 2.1: Example of an octree storing free (shaded white) and occupied (black) cells. The volumetric model is shown on the left and the corresponding tree representation on the right.

This subdivision continues until voxels are homogeneous or a minimum resolution/maximum depth is reached. This process creates a tree where large, uniform regions (whether free, occupied, or unknown) are represented by single nodes at a shallower depth, making it highly efficient for representing sparse environments in comparison to uniform grid voxelization. The **OctoMap** framework [18] uses this principle to build probabilistic 3D occupancy maps from noisy sensor data. In such a framework, unobservable space is typically represented implicitly by the absence of nodes in a given volume; querying such a region returns an “unknown” status. The same octree structure also underpins learning-based 3D methods such as OctNet

[27], which runs deep convolutions efficiently on an octree rather than building an occupancy map; here it is the occupancy-mapping use that is relevant.

Sidd’s Critical Regions. To make planning more efficient, space can be divided based on task-relevant meaning rather than uniform geometric criteria. One such approach is the concept of **critical regions (CRs)**, introduced by Molina et al. [24] and later extended by Shah and Srivastava [29]. In the original formulation, a critical region is a part of the configuration space with a low probability of being sampled by a uniform sampler yet essential to a solution—narrow passages such as doorways [24]. Shah and Srivastava generalised this to regions where task-relevant features—controller success rates, object visibility, or transition likelihoods—change significantly [29].

In this line of work the regions are not hand-specified: a deep network is trained to *predict* the critical regions of a given environment from that scene’s input, and the predictor generalises to environments it was not trained on. Once predicted, a region-based Voronoi diagram (RBVD) partitions the configuration space around the CRs: each point is assigned to the nearest critical region, forming abstract states. Transitions between these abstract states define abstract actions, forming the backbone for hierarchical planning. By focusing computation on these task-relevant regions, this approach scales planning more effectively, but the regions are predicted once per environment and then held fixed; they do not adapt as objects move within a scene during execution.

Our work combines these techniques, as explained in Chapter 4.

3 Related Work

Integrating task and motion planning with partial observability is challenging, particularly when a robot must act on incomplete information about its environment. This section reviews previous work on this topic, focusing on planning approaches and methods for representing a robot’s beliefs about the world.

3.1 TAMP under Partial Observability

Existing approaches to TAMP under partial observability generally either use formal, probabilistic models or more ad-hoc, reactive methods, or a combination of both.

3.1.1 POMDP-based TAMP Solutions

Several works formulate planning under uncertainty as a POMDP to handle it in a principled, risk-aware way [7, 28]. TAMP with Uncertainty and Risk Awareness (TAMPURA) [7], for example, learns a coarse abstract model of each given controller’s preconditions and effects, builds a non-deterministic MDP, and solves it with loop AO* search (LAO*), an uncertainty-aware solver [16], to produce risk-aware, information-gathering plans rather than provably optimal ones. Saleem et al. [28], by contrast, address contact-rich *manipulation* under object-pose uncertainty—not full task-and-motion planning—computing partial policies with an anytime heuristic-search POMDP solver. However, a persistent challenge is the representation of belief over continuous properties, particularly object poses. TAMPURA [7] samples object placements in continuous pose space and, in its real-robot pipeline, draws its object-pose belief from a perception front-end, Bayes3D [14]: a generative inverse-graphics model that infers a posterior over 3D scenes by rendering candidate object arrangements and scoring them against the observed depth image.

3.1.2 Partially Grounded Plans in TAMP

In contrast to formal probabilistic models, other TAMP systems handle uncertainty by deferring it to execution [26]. These methods generate partially grounded plans, leaving “gaps” for actions with unknown outcomes. During execution, specialized closed-loop behaviors (hand-designed or learned) attempt to fill these gaps. If a controller fails, the failure is treated as a new constraint, and the system replans. This reactive approach avoids the need to model outcome probabilities explicitly. However, because it does not reason about uncertainty proactively, it can struggle

with problems that require multi-step, strategic information gathering, such as deciding which of several actions is most likely to reveal a hidden object.

3.1.3 Modern Approaches to TAMP under Uncertainty

Recent works have leveraged learning and large models to address partial observability, offering different strategies than our own.

Ma et al. [22] add a strategic reasoning layer on top of a standard planner: they first reason backward from an underspecified goal to decide which objects matter (Task-level Backward Search), then use a constraint graph to order object moves collision-free through clutter (Object Manipulation Constraint Graph). We differ at the representation level rather than by adding a layer: we introduce a geometric belief representation, Boxels, so that the planner treats which regions are observed versus occluded as first-class planning state and information-gathering becomes a core planner action rather than a separate pre-planning step; the visibility signal is currently supplied by an oracle.

Another approach sources the planner’s belief from large foundation models. Zhao et al. [30] use a vision-language model (VLM) as an uncertainty estimator: the VLM answers the planner’s perceptual queries about an image (for example, “Is the drawer empty?”), and the resulting belief feeds a *symbolic* belief-space planner—a pairing of learned perception with structured reasoning that is reported to outperform end-to-end VLM planners. CoCo-TAMP [21] is similar in spirit, drawing on the commonsense priors of a large language model (LLM) to shape the belief over where task-relevant objects are likely to be, again within a partially observable TAMP loop. The difference from our work is the *source* of the belief: theirs is learned (a VLM or LLM), whereas ours is a structured, object-centric geometric partition into discrete semantic Boxels, with object poses currently supplied by a ground-truth oracle rather than a learned detector. Integrating a real detector is left to future work.

Finally, some methods move away from symbolic planning altogether and instead use end-to-end learning. Bai et al. [2] (TAVP), for example, use a separate reinforcement-learning policy to select task-relevant *virtual* camera viewpoints and re-render the scene from them—rather than physically moving the camera—and then produce the manipulation action with an end-to-end learned policy rather than a planner. This is fundamentally different from our model-based approach, where looking around is not a learned reflex but an explicit action chosen by the planner because its world model indicates that knowledge is missing. Because our method separates planning from execution, it is more transparent and can be retargeted to new goals without retraining a policy; the cost, however, is that a genuinely new goal requires authoring new PDDL actions and their streams, not just a new goal specification.

3.1.4 Belief-Space Planning and Replanning

A long line of work plans directly in *belief space*, treating the robot’s distribution over world states as the object of planning. Kaelbling and Lozano-Pérez [20] plan over the robot’s belief using hierarchical goal regression, while Hadfield-Menell et al. [15] introduced a modular interface that exchanges information between an off-the-shelf classical planner and the geometric layer as the plan is built, determining the problem under the maximum-likelihood-observation assumption so that a successful plan generates the observations it needs. A practical alternative to maintaining an explicit stochastic model is to *determinise and replan*: SS-Replan [11] plans deterministically in a hybrid belief space and replans whenever a new observation arrives, including sensing actions within a single plan. Our system shares this determinise-and-replan strategy—we commit to an optimistic deterministic plan and replan on each observation rather than synthesising a probabilistic policy. What we add is separate from the replanning loop: a spatial belief in which occlusion is explicit, named planning state, so information-gathering is explicitly composed within the planning representation.

3.1.5 Partially Observable Deterministic Planning

A separate tradition handles partial observability *without* probabilities, by reasoning about what is *known*. A prominent line within partially observable deterministic (POD) planning represents knowledge with explicit *knowledge literals* and reduces the problem to classical planning by compilation; other POD planners instead search belief space directly, as in Contingent-FF and MBP. The translation-based contingent planner CLG [1] and the classical-replanning K-replanner [4] established this reduction; LW1 [5] made it scalable with a *linear* translation; and a parallel route compiles partial observability to fully-observable non-deterministic (FOND) planning [25]. The underlying models are treated by Geffner and Bonet [12]. This thesis builds directly on the POD line: we carry knowledge literals—e.g. whether a region’s contents are known—inside the PDDL state and let a classical planner reason over them.

To our knowledge, the deterministic, knowledge-literal POD line has not previously been combined with task and motion planning, nor used a geometric belief representation that makes *which regions are occluded* part of the symbolic state. Belief-space TAMP systems do reason about occlusion symbolically—IBSP carries maximum-likelihood occlusion fluents [15] and SS-Replan derives an occlusion predicate by ray-casting over its pose belief [11]—but in both the occlusion is a check over a *probabilistic* belief about object poses, not a named occluded *volume* held as deterministic symbolic state. Occlusion-aware object retrieval has likewise been studied—for instance by searching the likely locations of an occluded target drawn from a *learned* distribution [3]—and partially observable TAMP systems increasingly source their belief from learned models [21, 30] or probabilistic abstractions [7]; but we are not aware of prior work that represents occluded volumes as named, object-centric symbolic state resolved by sensing within a PDDLStream TAMP loop—the gap this thesis addresses.

3.1.6 Summary and Research Gap

These approaches occupy distinct points in the design space. POMDP-based methods such as TAMPURA [7] reason about uncertainty proactively but rely on a learned probabilistic transition model, and they keep the spatial belief sub-symbolic, so occluded regions are never named, inspectable state the planner can target. Reactive, partially-grounded methods [26] avoid probabilistic modelling but cannot plan multi-step information-gathering. Belief-space replanning [11] determinises and replans, as we do, but does not make occlusion a first-class planning entity. The POD tradition [5] reasons about knowledge symbolically and deterministically, but has not been applied to TAMP or to spatial occlusion. We build on these lines: a POD task-and-motion planner whose spatial belief is an adaptive, object-centric partition (Boxels) in which occluded regions and their observed/occluded status are explicit symbolic state. The planner thus reasons about *where* information is missing and composes deliberate look-and-clear actions, without the cost and opacity of a learned probabilistic model.

This position yields three advantages, each with a cost we state plainly. (i) *Occlusion is inspectable state*: named observed-versus-occluded regions a reader can enumerate, rather than a sub-symbolic visibility grid feeding a learned outcome model [7]; the cost is that the visibility signal is currently supplied by an oracle rather than a perception module. (ii) *Uncertainty is handled without probabilities*: knowledge literals record what is known deterministically, with no per-action probability to estimate; the cost is that the planner cannot rank actions by success likelihood and instead minimises plan cost and replans when an optimistic assumption fails. (iii) *No MDP machinery is required*: a lightweight deterministic PDDL layer replaces the learned-MDP stack of POMDP-based TAMP [7]; the cost is that optimistic determinise-and-replan is sound only while belief reflects observation, an invariant the bounded give-up on unreachable shadows (Section 6.5) deliberately relaxes.

3.2 Spatial Belief Representation in TAMP

In a typical TAMP cycle the planner (i) reasons symbolically over high-level actions (pick, place, move-to-region), (ii) instantiates them with continuous parameters using a motion planner, (iii) executes the resulting trajectory, observes the environment, and (iv) updates its belief before the next planning step. Every one of these phases depends critically on how we represent spatial belief. While general methods were introduced in Section 2.4, their application in TAMP reveals important limitations.

3.2.1 Limitations of Geometric Discretization for POD-TAMP

Octrees: The efficiency of Octrees makes them a particularly useful tool for 3D mapping. While an Octree framework like OctoMap [18] can represent unobserved space, its primary

purpose is to efficiently model occupancy based on noisy sensor measurements. For the task of planning to find hidden objects, it has two key limitations:

1. **Lack of Semantic Structure:** Standard Octree subdivision is purely geometric and based on occupancy variance. It does not create semantically meaningful regions corresponding to, for example, "the volume occluded by object A". The planner would have to infer this semantic meaning from the raw geometric data, which is a complex task in itself.
2. **Object Integrity:** The recursive subdivision of an Octree can arbitrarily split a single physical object into multiple voxels at different tree depths. This complicates reasoning about objects as whole entities, which is essential for manipulation planning (e.g., "pick up object A," not "pick up voxels 1, 2, and 5 of object A").

3.2.2 Limitations of Static Semantic Abstractions

Sidd's Critical Regions (CRs): The Critical Regions method [29] uses a learned predictor to identify task-relevant regions for each environment from that scene's input, giving the planner a meaningful, non-uniform way to abstract the space. The predictor generalises across environments, but within a given scene the regions are predicted once and then held fixed, so the partition does not adapt when objects move during execution. For example, the region "behind object A" depends on where A currently is and where the robot is standing—a partition fixed at the start of the episode cannot track that.

4 Approach

This chapter presents our approach to bridging abstract belief reasoning and continuous geometry: a spatial abstraction called Boxels that partitions the robot’s workspace into meaningful regions. We integrate these into a POD TAMP framework through three components: (1) an adaptive discretization that builds Boxels from detected objects and their occlusions, (2) a formal model that represents belief about Boxel occupancy as symbolic facts, and (3) a PDDLStream implementation that grounds these facts in geometric and sensing operations.

4.1 Overview

Our methodology links abstract belief representation to continuous geometric reasoning through three components:

1. **Adaptive Semantic Discretization:** A dynamic procedure for partitioning the robot’s workspace into a set of meaningful Semantic Boxels, i.e. semantically meaningful subspaces, based on detected objects and their known 3D models.
2. **A Formal POD-TAMP Model:** A formal model that uses these Boxels as the foundation for its state representation, allowing a POD planner to reason about knowledge and uncertainty regarding object locations and unobstructed spaces.
3. **PDDLStream Integration:** An implementation using custom PDDL fluents, actions, and streams that operate on the belief state over Boxels, realizing the formal model in PDDLStream.

4.2 Adaptive Semantic Discretization: Generating Boxels

We build the Boxel abstraction through a process of *adaptive semantic discretization*. This process dynamically generates a set of Semantic Boxels ($\mathcal{BX}_{\text{set}}$) by partitioning the workspace based on detected objects and the occlusions they create. It runs as a Python stage before each call to the planner—and is re-run between actions as the scene changes—rather than as a PDDLStream procedure. Figure 4.1 illustrates the process.




Figure 4.1: Adaptive Semantic Discretization. (a) Initial scene: the camera viewpoint and the detected objects on the table, with a fully occluded target (cyan) sitting in the shadow behind the largest occluder. (b) Each object is bounded by a dedicated object Boxel. (c) The volume each object hides from the camera becomes an occlusion (shadow) Boxel, swept from the object's back face out to the workspace boundary along the lines of sight (dashed) and rising to the object's height. (d)–(f) The remaining free space is partitioned by an octree: from a single cell spanning the whole workspace (d), only cells containing both free and occupied space are recursively subdivided—in all three dimensions, so the open space above the objects stays a few large cells (e)—and adjacent free cells are then greedily merged across aligned faces into larger convex Free Space Boxels (f).





Figure 4.2: The same partition applied to a PyBullet scene. The cyan wireframe cells are the Free Space Boxels filling the workspace; at the cube cluster, object bounding cuboids enclose the visible cubes and red wireframe Boxels mark the volumes they occlude.


The generation process is as follows:

1. **Object-Centric Bounding:** When an object is detected, we generate a dedicated Boxel: a cuboid that tightly bounds the object's known 3D model at its detected pose. This isolates each known object in its own spatial region. Object detection is provided by an oracle that returns ground-truth object identities and poses from the simulator, which isolates the planning contribution from perception noise; replacing it with a learned, image-based detector is left to future work.
2. **Occlusion-Aware Subdivision:** For each object-bounding Boxel, the space occluded by it from the fixed camera viewpoint is calculated and partitioned into one or more shadow Boxels, split by any intervening visible obstacles. As the robot acts, the set of shadow Boxels is updated: sensing a region resolves its shadow, and discovering a previously hidden object adds a shadow for that object. A relocated object, however, is not given a new shadow at its destination. Re-creating shadows for moved objects could admit plans that never terminate, so it is deliberately omitted; handling it safely would require additional belief bookkeeping and is left to future work. This explicitly represents regions the robot currently cannot observe, so the planner can target them with information-gathering actions.
3. **Recursive Partitioning:** The space not occupied by object or occlusion Boxels is partitioned recursively, starting from a single Boxel that spans the whole workspace. A Boxel found to be empty is kept as a Free Space Boxel; a Boxel that still overlaps an object or occlusion is split into eight octants, and the procedure repeats on each. Subdivision stops when a region is empty or a minimum Boxel size is reached. The resulting free cells are then merged greedily across shared faces into larger convex Boxels, which reduces their number without changing the represented free volume. This merge is convex-only—two cells are combined only when they share an exactly aligned face—so the partition retains more, smaller free Boxels than a full semantic merge would; this is accepted as sufficient for the object counts evaluated here (see Section 6.5).

This adaptive, object-centric approach contrasts with a uniform voxel grid by focusing discretization effort on task-relevant areas—the objects themselves and the occluded spaces they create—rather than partitioning all space uniformly. Its free-space stage is itself an octree, and a uniform-grid baseline is implemented for comparison. Each generated Boxel is assigned a unique ID and associated geometric data; the completed set of Boxels is supplied to the planner as static facts in its initial state. Figure 4.3 contrasts the two partitions on a matching scene.



(a) Semantic partition: cyan Free Space Boxes where the scene is empty, plus object and shadow Boxels at the cube cluster.



(b) Uniform-grid baseline: hundreds of fixed-size cyan cells covering the whole workspace.

Figure 4.3: Adaptive semantic partition (left) versus uniform-grid baseline (right) on a matching scene. The adaptive partition produces a small set of task-relevant cells; the uniform grid covers the whole workspace at the same cell size, producing an order of magnitude more cells.

4.3 Belief over Boxels with Know-If Fluents

The Semantic Boxels generated in the previous section partition the workspace into a finite set of regions; what the robot does not know is which Boxel a hidden object occupies. We represent this uncertainty with the Know-If Fluents introduced in Chapter 2. For each object obj and Boxel $Boxel$, the belief carries the K-literal $K(\text{InBoxel}(obj, Boxel))$ when the robot knows the object is in that Boxel and $K(\neg \text{InBoxel}(obj, Boxel))$ when it knows the object is not; the absence of both means the object is possibly there. Because the set of Boxels is finite, the belief over object locations is finite as well, and it is held directly in the planner’s symbolic state.

Representing belief this way is what keeps planning tractable. As described in Chapter 2, the K-literal translation compiles a partially observable problem into a classical one. Applied here, it lets an ordinary classical planner reason directly about where objects might be and plan sensing actions to resolve that uncertainty within the same classical search, rather than delegating belief to a separate contingent or belief-space planner. This compile-to-classical reduction is the one introduced by the POD planners CLG and LW1 [1, 5]; we adopt their knowledge-literal translation and solve the result by optimistic determinisation and reactive replanning (Chapter 3): we assume each sensing action succeeds (optimistic), plan as if the world were fully known, execute, and replan whenever an observation contradicts that assumption (reactive replanning).

Knowing which Boxel holds an object is not enough on its own: a manipulation action also needs a reachable grasp and a collision-free trajectory. For this continuous-geometry side we rely on PDDLStream, which extends classical PDDL with *streams*—the lazily-evaluated

samplers and tests for continuous parameters described in Chapter 2. Our approach combines the two in a single PDDL domain: Know-If Fluents carry the discrete belief over Boxels, while streams certify the continuous facts—grasps, inverse-kinematics solutions, and motions—that the manipulation actions require. A symbolic action such as picking a target object therefore mixes both kinds of precondition: a Know-If Fluent stating that the target’s Boxel is known, and stream-certified geometric facts stating that the grasp and motion are feasible. The following two sections make this concrete, giving first the formal POD-TAMP model and then its PDDLStream realization.

4.4 Semantic POD-TAMP Model Definition

We formalize our Semantic POD-TAMP system as an explicit state model, adapting the state-space formulation of Geffner and Bonet [12] to the partially observable deterministic setting. The model is a tuple:

$$\mathcal{M} = \langle \mathcal{S}, \mathcal{S}_0, \mathcal{S}_G, \mathcal{A}, f, \mathcal{O} \rangle$$

where:

- **$\mathcal{S}$ (States):** A finite set of abstract states. Each state is a consistent assignment of truth values to a set of atomic propositions $\mathcal{AP}$ built from K-literals (see Chapter 2). These propositions are grounded in the Boxel representation, e.g., $K(\text{InBoxel}(\text{obj}_k, \text{Boxel}_j))$.
- **$\mathcal{S}_0$ (Initial States):** The set of possible initial abstract states, representing the initial belief b_0 .
- **$\mathcal{S}_G$ (Goal States):** The set of goal states to be reached with certainty, specified by a logical formula over $\mathcal{AP}$ (e.g., $K(\text{holding}(\text{target_obj}))$).
- **$\mathcal{A}$ (Actions):** A set of abstract actions defined as PDDLStream action schemas that operate on the belief state, such as ‘sense’ or ‘pick’.
- **f (Transition Function):** A deterministic state-transition function $s' = f(a, s)$ that applies an action’s add and delete effects to the current state, so that the successor state reflects the updated knowledge.
- **$\mathcal{O}$ (Sensor Model):** The sensor model relates a sensing action to the observation it yields—the target is found in a Boxel, or it is not—and the resulting belief update. In our system this is not realized as an observation function inside the POD planner: the planner’s sensing action is optimistic and assumes the target is found, while the true outcome is obtained at execution time and folded into the belief by a Python sense-plan-act loop, which replans whenever the optimistic assumption proves wrong.

The planning problem is to find a sequence of actions from $\mathcal{A}$ that moves the belief state from $b_0 = \mathcal{S}_0$ to a belief state b_g in which every possible state is a goal state (i.e., $b_g \subseteq \mathcal{S}_G$).

4.5 PDDLStream Integration

This section presents the implemented PDDL domain and streams that realize the model on the belief state over Boxels. For readability the listings show the belief-related fragment of the domain, eliding the geometric and stacking predicates that are not central to this discussion.

4.5.1 PDDL Fluents for Belief Representation

The agent's knowledge is represented directly in the PDDL state. The implemented domain is untyped STRIPS, so object categories are themselves predicates, and the belief over object locations is carried by a single Know-If fluent.

```

1  (:predicates
2    ; --- Category predicates (untyped domain) ---
3    (Boxel ?x)           ; ?x is a Boxel
4    (Obj ?x)             ; ?x is a manipulable object
5
6    ; --- Boxel classification ---
7    (is_shadow ?b)       ; ?b is an occluded region
8    (is_object ?b)       ; ?b bounds a physical object
9    (is_free_space ?b)   ; ?b is known-empty space
10
11   ; --- Object location: value and knowledge ---
12   (obj_at_boxel ?o ?b)   ; object ?o is at Boxel ?b
13   (obj_at_boxel_KIF ?o ?b) ; Know-If fluent: it is known whether ?o is at ?b
14
15   ; --- Other knowledge and robot state ---
16   (obj_pose_known ?o)    ; ?o has been localized to a Boxel
17   (handempty)            ; the gripper is empty
18   (holding ?o)           ; the gripper holds ?o
19   ; ... further robot, geometry, and stacking predicates
20 )
```

Listing 4.1: PDDL fluents for belief over Boxels

Section 4.3 describes belief with the two K-literals $K(\text{InBoxel})$ and $K(\neg \text{InBoxel})$. The implemented domain realizes the same belief with a single Know-If fluent: `obj_at_boxel_KIF` marks that the location of an object relative to a Boxel is *known*, while `obj_at_boxel` carries the value; the absence of the Know-If fluent means the location is still uncertain. The two encodings are logically equivalent, and the single fluent keeps the grounded state smaller.

4.5.2 The Sensing Action

Information gathering is realized by a single sense action, which checks whether a target object is present in a given Boxel.

```

1  (:action sense
```

```

2      :parameters (?o ?region)
3      :precondition (and
4          (Obj ?o)
5          (Boxel ?region)
6          (view_clear ?region)           ; line of sight to the region is clear
7          (boxel_fits ?o ?region)        ; ?o physically fits inside ?region
8          (not (obj_at_boxel_KIF ?o ?region)) ; only sense while the location is
              unknown
9      )
10     :effect (and
11         (obj_at_boxel_KIF ?o ?region)    ; the location of ?o is now known
12         (obj_at_boxel ?o ?region)        ; optimistic: assume the object is
              found here
13         (obj_pose_known ?o)
14         (increase (total-cost) 1)
15     )
16 )

```

Listing 4.2: The implemented sense action

The action is *optimistic*: its effect unconditionally asserts that the object is found in the sensed region. An earlier design used conditional effects that branched on a found/not_found observation, which would let the planner build a multi-step search (“sense *A*; if empty, sense *B*”) within one plan. PDDLStream and its classical backend do not support such contingent planning, so the implemented domain commits to the optimistic outcome and relies on *reactive replanning*: when execution reveals that the target is not in the sensed region, the execution loop updates the belief and replans. An empty region is recorded as clear; a region that instead turns out to hold a *non-target* object has that object—and the new region it occludes—added to the partition, re-boxelizing the workspace before the next plan so the newly revealed occluder enters the belief rather than being discarded. With N candidate regions this converges in at most N replanning cycles and is functionally equivalent to a conditional plan executed in sequence. Figure 4.4 shows the action targeting a shadow region in the simulator.




Figure 4.4: The sense action targeting a shadow region in PyBullet. The arm has been moved to a sensing configuration over a labelled shadow; the in-simulator overlay labels each object and its shadow region. This is the configuration in which the sense action’s (`view_clear ?region`) precondition is checked and its optimistic (`obj_at_boxel ?o ?region`) effect is asserted.

One bounded exception to this observation-driven belief update concerns regions that cannot be observed at all: if the line of sight to a shadow remains blocked after a fixed number of sensing attempts, that shadow is marked empty without ever being directly observed. This keeps the planner progressing when a region is effectively unreachable, but it is unsound—the region could still hold the target. A sound treatment, which would re-ground the view-blocking facts so the planner first clears the obstruction, is left to future work.

The sense action calls no streams. Because a fixed overhead camera observes the scene, no robot sensing configuration has to be sampled, and the only geometric precondition is the derived predicate (`view_clear ?region`), which holds when no object blocks the line of sight to the region. Streams instead serve the manipulation actions: the domain declares `sample-grasp`, `plan-motion`, `compute-kin`, and `compute-stack-kin`, so that an action such as `pick` requires a valid grasp, an inverse-kinematics solution, and a collision-free motion, each certified by a stream.

5 Evaluation

This chapter reports the experimental evaluation of the Boxel-based POD-TAMP framework. Section 5.1 describes the experimental setup as executed. Section 5.2 defines the measured quantities. Section 5.3 defines the variants compared, including the TAMPURA external reference. The results themselves follow in the remaining sections.

5.1 Experimental Setup

Simulator and robot. All experiments were run inside PyBullet [6], with a 7-DOF Franka Emika Panda as the manipulator on a tabletop.

Goals. Three goals were evaluated (Figure 5.1):

- **holding:** the robot must end the episode holding a single target object whose initial pose is hidden behind one of a set of occluders. The number of occluders n_{occ} is the difficulty axis (2, 3, 4; see Figure 5.2).
- **stack:** the robot must build a stack of cubes to a given height. The stack height h is the difficulty axis (2, 3, 4); cube positions are fully observable.
- **find-and-tray-stack:** the robot must locate hidden cubes behind occluders and stack them on a tray. The number of occluders n_{occ} is the difficulty axis (2, 3, 4).




Figure 5.1: The three evaluated goals. **holding:** retrieve a single target hidden behind occluders. **stack:** build a tower of cubes to a target height. **find-and-tray-stack:** locate hidden cubes and stack them on a tray.




Figure 5.2: The occluder-count difficulty axis n_{occ} , shown at low, medium, and high settings. More occluders crowd the target and occlude a larger share of the workspace, so the planner must reason over more shadow Boxels.

Scenes and seeds. Each (goal, variant, difficulty) combination was run on 100 random scenes; seeds are paired across the `semantic` and `semantic+mbs0.05` variants on the random-pairs scene, so their head-to-head comparison is seed-for-seed. The headline comparison spans ≈ 2700 cells ($3 \text{ goals} \times 3 \text{ variants} \times 3 \text{ difficulty levels} \times 100 \text{ seeds}$, with minor rounding from scheduling); together with the coarser free-space resolution-floor arms swept in Section 5.4.5 ($1.5\times$ and $2\times$ the auto leaf size), a total of ≈ 4500 cells were executed.

Perception. Object detection is performed by an oracle that reads ground-truth object poses from the simulator and uses ray-casts from a fixed overhead camera (Figure 5.3) to determine which objects, and which Boxels, are visible. The simulator also renders RGB-D images, but they are not processed inside the planning loop. The modelled uncertainty is occlusion uncertainty—which Boxel a hidden object occupies—resolved at run time by the sense action.




Figure 5.3: An evaluation scene observed by the fixed overhead camera, shown here from an oblique simulator viewpoint; the corner panels are the camera’s synthetic RGB (top) and depth (bottom) views. Object detection is an oracle that reads ground-truth object poses from the simulator rather than processing the rendered RGB-D images. A future learned detector would consume those images in place of the oracle, without changing the planning architecture.

Planning budget. Each PDDLStream call (initial plan or replan) was given a wall-clock budget of 1800 s (30 min); a cell terminates with the `timeout` exit reason if it exhausts this budget. Per-call planning time varies by goal: over successful cells the median call takes 0.97 s for stack, 2.05 s for find, and 2.97 s for find and stack; the 1800 s budget is a safety bound hit only by rare pathological or timeout cells, not the operating point. The anytime sweep records cumulative solve rate as a function of wall-clock budget, and the curves are essentially flat past ≈ 200 s.

Software. PyBullet [6]; PDDLStream [10] with the FastDownward classical planner backend; Python 3.10.

5.2 Evaluation Metrics

For each cell we record:

- **Success:** whether the robot reached the goal within the planning budget, without exceeding the per-episode replan limit.
- **Planning time:** total wall-clock time spent inside PDDLStream calls, summed across replans. Wall-clock per cell is effectively equal to planning time; execution, perception, and motion overhead are negligible at the scales evaluated.

- **Replan count:** how many times the planner was re-invoked in the sense-plan-act loop before terminating.
- **Representation size:** the Boxel set decomposed into OBJECT, SHADOW, and FREE_SPACE cells, and the number of grounded facts in PDDLStream’s initial state. These two capture the per-cell cost of the partitioning strategy independently of planner choice.
- **Failure mode:** when a cell does not succeed, an exit reason among `planner_failed`, `timeout`, `replan_limit`, `physics_mismatch`, `drop_failed` (the gripper cannot release a non-target object; the episode exits after three failed drop events), `no_summary`, and `all_searched`.

Aggregate means in this chapter are taken either over all n_{cells} or, for quantities that are only defined on cells that produced a snapshot (notably mean initial-state facts and mean replan count), over the subset of cells with non-null values. These two denominators differ by about 6 % of cells—those that timed out early, before any planner state was recorded.

5.3 Baselines for Comparison

Three variants of the planner are compared, plus one external reference:

1. **semantic** — the adaptive Boxel partition described in Chapter 4. Free-space cells are placed only where the workspace is known to be empty; their leaf size is auto-set to the largest object’s footprint ($\approx 6\text{--}9$ cm). This floor is a deliberate design choice rather than a tuning limit: each object is bounded by exactly one Boxel, so leaves finer than an object’s footprint would fragment placement targets on the table dimensions used here.
2. **semantic+mbs0.05** — the same adaptive partition but with the free-space leaf-size floor forced to 5 cm (`min_boxel_size = 0.05`). Because the auto-set leaf size is already coarser ($\approx 6\text{--}9$ cm) and the convex merge re-absorbs finer free cells, this floor never binds: the partition is identical to plain **semantic** (matching to $\sim 0.1\%$). The variant therefore serves as a check that a finer free-space floor does not change the partition, not as a separate operating point.
3. **uniform** — the adaptive partition is replaced by a uniform voxel grid over the workspace; the object and shadow Boxels are unchanged. The cell size is auto-set to the same largest-object-footprint value, since finer cells likewise break placement. This isolates the contribution of the adaptive free-space partition.
4. **TAMPURA** [7] — a published POMDP-based TAMP framework. Its Table II reports several model-learning/search configurations on a *Partial Observability* task; we compare against TAMPURA’s own primary configuration there—Bayes-Optimistic belief updates with LAO* search—which it reports at 57 ± 38 s over 20 trials (Table II of [7], arXiv v2), rather than re-running it; the comparison is architectural — a probabilistic learned-MDP

planner solved with LAO* versus our deterministic knowledge-literal planning with replanning — not a hardware benchmark. See Section 6.3 for the full framing.

5.4 Results

Table 5.1 reports the aggregate success rate and mean success-only planning time for each (goal, variant) cell. The semantic and uniform variants differ by an order of magnitude or more in both axes on every goal; semantic+mbs0.05 sits within seed noise of semantic on stack, modestly ahead on holding, and modestly behind on find-and-tray-stack.

goal	variant	n_{cells}	success rate	mean plan time (s)
find-and-tray-stack	semantic	299	39.8 %	28.37
find-and-tray-stack	semantic+mbs0.05	300	33.0 %	21.49
find-and-tray-stack	uniform	300	1.3 %	156.10
holding	semantic	300	42.3 %	7.78
holding	semantic+mbs0.05	300	46.3 %	7.42
holding	uniform	300	33.3 %	50.35
stack	semantic	300	61.3 %	1.23
stack	semantic+mbs0.05	300	61.3 %	1.23
stack	uniform	301	39.2 %	23.84

Table 5.1: Headline success rate and mean planning time per (goal, variant). Planning time is success-only.

5.4.1 Anytime Solve Rate

Figure 5.4 plots cumulative solve rate against wall-clock budget on a log axis. Three observations stand out. First, the semantic curves rise to their final value within ~ 10 – 30 s on every goal, and the additional budget out to 1800 s buys at most a few extra successes; the operating point is the first few minutes, not the long tail. Second, on stack the semantic and semantic+mbs0.05 curves overlap exactly — the 5 cm leaf floor changes nothing here. Third, the uniform curve is delayed in onset on every goal and tops out lower; on find-and-tray-stack it never gets off the ground (4 solves out of 300).




Figure 5.4: Cumulative solve rate vs wall-clock budget (log axis) per (goal, variant). Almost all solves happen within ~ 200 s; uniform starts later and plateaus lower on every goal.

5.4.2 Success Rate and Scene Difficulty

Figures 5.5 to 5.7 show success rate against the difficulty axis for each goal. On holding (Figure 5.5), success is roughly flat in $n_{\text{occ}} \in \{2, 3, 4\}$: semantic stays at 39–47 %, semantic+mbs0.05 at 44–48 %, uniform at 31–36 %. There is no clear scalability degradation in this range, only a ~ 10 pp gap between adaptive and uniform partitioning.




Figure 5.5: Success rate vs n_{occ} on the holding goal. Adaptive variants (semantic, semantic+mbs0.05) are within seed noise of each other; uniform sits ~ 10 pp lower. Each bar is the mean success rate over the ~ 80 – 100 seeds at that difficulty (here and in Figures 5.6 and 5.7).

On stack (Figure 5.6), the difficulty axis is the stack height. All variants solve $h = 2$ trivially (97 %); semantic and semantic+mbs0.05 degrade gracefully to 74 % at $h = 3$ and 13 % at $h = 4$,

while uniform collapses from 97 % to 20 % to 1 %. The semantic and semantic+mbs0.05 curves are indistinguishable seed-for-seed.




Figure 5.6: Success rate vs stack height h on the stack goal. Adaptive variants degrade gracefully; uniform collapses past $h = 2$.

Find-and-tray-stack (Figure 5.7) is the discriminating goal: uniform succeeds on at most 2 out of 100 cells at every difficulty, whereas semantic stays flat around 40 %. The semantic+mbs0.05 variant degrades with n_{occ} (40 % $\rightarrow$ 31 % $\rightarrow$ 28 %) on this goal, the only case in the sweep where a finer leaf floor visibly hurts.




Figure 5.7: Success rate vs n_{occ} on the find-and-tray-stack goal. Uniform fails almost entirely; semantic+mbs0.05 visibly degrades with n_{occ} , contrary to its intent.

5.4.3 Planning Time and Scene Difficulty

Figure 5.8 shows mean success-only planning time on holding against n_{occ} . The adaptive variants stay below 20s across the range; uniform climbs from 21 s at $n_{\text{occ}} = 2$ to 83 s at $n_{\text{occ}} = 4$, four to five times the adaptive cost at each point. Planning times on stack and find-and-tray-stack follow the same pattern (semantic 1–2 s on stack and ~ 20 –50 s on FATS; uniform 20–220 s and 25–533 s respectively).




Figure 5.8: Mean success-only planning time vs n_{occ} on the holding goal. Adaptive variants scale gracefully; uniform’s cost grows much faster.

5.4.4 Representation Compactness

The cost gap traces back to the size of the representation. Figure 5.9 stacks the OBJECT, SHADOW, and FREE_SPACE Boxel counts per variant on holding. The adaptive variants produce ~ 25 –43 Boxels total; the uniform variant produces ~ 300 free-space cells alone. On the stack goal the ratio is steeper (semantic ~ 25 vs uniform ~ 1340), and the grounded-fact count scales with it: ~ 300 facts for the adaptive variants vs 11 000–14 000 for uniform across stack heights. Larger init-state $\Rightarrow$ larger PDDL grounding $\Rightarrow$ longer per-call planning time, which is the mechanism behind the timing gap in Figure 5.8.




Figure 5.9: Mean Boxel count by category, holding goal. Adaptive variants produce an order of magnitude fewer cells than the uniform baseline.

5.4.5 Free-Space Resolution

The adaptive partition has one free parameter: the free-space leaf size. The semantic variant auto-sets it to the largest visible object’s footprint plus a 1 cm margin ($\approx 6\text{--}9$ cm, scene-dependent). To check whether this heuristic sits at a sensitive operating point, we swept the leaf size both below auto (the semantic+mbs0.05 arm, 5 cm) and coarser than auto, at $1.5\times$ and $2\times$ the auto value, on the same 100-seed scenes. Figure 5.10 plots the resulting Boxel count against the leaf size as grouped bars per goal.




Figure 5.10: Mean Boxel count vs free-space leaf size for the find-and-tray-stack and holding goals (semantic variant; total and free-space cells). Each bar group is one resolution arm, labelled by its mean effective leaf size. The count is flat at and below the auto-set leaf size and falls 30–40 % as the leaf is coarsened above it.

The partition shrinks as the leaf is coarsened above auto — the free-space Boxel count falls by 30–40 % from $1\times$ to $2\times$ (holding $35 \rightarrow 20$, find-and-tray-stack $45 \rightarrow 27$, stack $28 \rightarrow 22$), with the grounded-fact count falling in step. Below and at auto the count is flat: forcing the leaf floor

to 5 cm leaves the partition unchanged (and on stack, identical seed-for-seed; Section 5.4.1), because the auto value already sits at or above the floor. Success rate is unmoved across the whole range: within seed noise on holding (42.3 % / 40.3 % / 45.7 % at $1\times$ / $1.5\times$ / $2\times$) and find-and-tray-stack (39.8 % / 35.3 % / 35.3 %), and identical seed-for-seed on stack (61.3 % throughout). Across roughly a fourfold range of leaf sizes ($\approx 5\text{--}18\text{ cm}$), then, the partition is insensitive to the setting: the auto-set heuristic sits on a flat plateau rather than at a tuned optimum, so coarsening buys a more compact representation at no measured success cost. Chapter 6 takes up why the partition is so insensitive.

5.4.6 Comparison with TAMPURA

Figure 5.11 compares per-episode times on our holding task — the closest analogue to TAMPURA’s hidden-object *Find Die* task — against real TAMPURA, re-run on the same hardware. TAMPURA’s published number for this task ($57.0 \pm 38\text{ s}$, the *Partial Observability* column of Table II in [7]) is a *per-step* planning time that includes executing the selected controller in simulation, not a per-episode figure; to obtain a like-for-like quantity we re-ran TAMPURA’s released *Find Die* environment on this thesis’s machine (20 seeds). Per successful episode it takes $166 \pm 85\text{ s}$, against our semantic variant’s 13.7 s wall-clock (success-only, median 8.5 s) — both measured end-to-end, including simulated execution and replanning, on the same CPU. Against the published per-step number the picture is unchanged: our matched per-step cost (wall-clock divided by the average number of plans) is 5.9 s . Reliability runs the other way — our semantic variant succeeds on 42 % of holding episodes, while the local TAMPURA run succeeds on 55 % (consistent with the $\geq 63\%$ implied by its 0.63 ± 0.30 discounted return, Table I, $\gamma = 0.98$, binary terminal reward). We are roughly an order of magnitude cheaper but less reliable, on a marginally easier goal (*Find Die* also requires returning home; our holding goal does not). The comparison is architectural — a probabilistic learned-MDP planner (LAO*) versus our deterministic knowledge-literal planning with replanning — and the two run on different environments, so it benchmarks neither system; Chapter 6 unpacks the framing.




Figure 5.11: Per-episode time on holding (same hardware): our wall-clock (semantic variant; mean $\pm$ std, success-only; right-skewed, median 8.5 s) and real TAMPURA *Find Die* re-run locally (mean $\pm$ std over successful episodes; 20 seeds, 55 % success). Both are per-episode times including simulated execution. TAMPURA’s published 57 s (Table II) is a per-step figure and is not plotted; see the text for the per-step reconciliation. On reliability the order reverses, holding-semantic succeeding on 42 % of episodes against TAMPURA’s local 55 %.

5.4.7 Failure Modes

Figure 5.12 stacks the exit reasons across all 6 (goal, variant) cells. `planner_failed` is the dominant failure mode in every condition; `replan_limit` appears predominantly on stack; `timeout` accounts for a small slice of find-and-tray-stack failures. On find-and-tray-stack the uniform bar is almost entirely `planner_failed`, consistent with the 1 % success rate.




Figure 5.12: Stacked exit-reason breakdown per (goal, variant). `planner_failed` here also absorbs the `physics_mismatch`, `drop_failed` and `all_searched` exits (all cases where the planned solution ultimately failed); it is the dominant failure mode on holding and find-and-tray-stack, while on stack the semantic variant instead fails mainly via `replan_limit` (the second-largest slice after success).

The qualitative readings of these figures, together with their architectural implications and threats to validity, are interpreted in Chapter 6.

6 Discussion

The evaluation in Chapter 5 measures the Boxel-based POD-TAMP framework on three goals (holding, stack, find-and-tray-stack) against a uniform-grid baseline that keeps everything else fixed, and against TAMPURA’s published numbers as an external reference. This chapter interprets those results: where the adaptive partition pays off, where finer leaves stop helping, and what the TAMPURA comparison does and does not show.

6.1 Adaptive vs Uniform Partitioning

The headline finding is the cost gap between the adaptive partition and a uniform free-space grid at the same cell scale. Both partitions use the same auto-set cell size — the largest object’s footprint, $\approx 6\text{--}9$ cm in our scenes — so the comparison isolates the structural choice rather than the resolution. On holding, the adaptive partition produces 25–43 Boxels per scene; the uniform grid produces ~ 320 free-space cells alone (Figure 5.9). Larger partitions translate to larger initial-state fact counts (~ 300 vs 11 000–14 000 on stack) and longer per-call planning time, which is the mechanism behind Figure 5.8.

What that gap buys is goal-dependent. On holding the uniform baseline still solves about a third of scenes; it pays several times the planning time at every n_{occ} but the mechanism still works. On stack the adaptive variants degrade gracefully from 97 % at $h = 2$ to 13 % at $h = 4$, while uniform collapses from 97 % to 1 %. On find-and-tray-stack, which composes occluded sensing with stacking, uniform is effectively broken (1.3 % success). The reading is that the adaptive partition does not change the planner; it changes how many irrelevant cells the planner has to ground and search over. When the goal needs only a handful of candidate placements — a single hidden target — the overhead of a few hundred free-space cells is tolerable; when the goal requires composing placements across multiple replans, the overhead dominates.

6.2 The 5 cm Leaf-Floor Variant

The `semantic+mbs0.05` variant was included to test whether forcing a 5 cm floor on the free-space leaf size changes outcomes against the default `semantic` configuration, where the leaf floor is the largest object’s footprint. It does not: on stack the curves overlap seed-for-seed; on holding the floor is ~ 4 pp ahead of `semantic`; on find-and-tray-stack it is ~ 7 pp behind, with a visible degradation as n_{occ} grows (Figure 5.7). The mean replan count rises slightly on holding

and find-and-tray-stack relative to semantic — consistent with more sensing/replanning cycles without more successes.

The mechanism is that the largest object’s footprint ($\sim 6\text{--}9$ cm in our object set) is already the binding constraint on placement; the 5 cm floor cannot shrink cells past this point. A finer free-space floor has little effect on cell count in any case—the greedy convex merge re-absorbs the smaller leaves, leaving the count essentially unchanged (only a coarser floor reduces it). In scenes whose placements are constrained by sub-5 cm geometry — narrow pockets, dense small-occluder shelves — a finer leaf floor would in principle help. Constructing such a scene and re-running the same comparison is the natural follow-up; it is noted in Section 7.1.

6.3 Architectural Comparison with TAMPURA

Figure 5.11 sets our holding results against real TAMPURA, re-run on the same hardware; holding — find a hidden object and pick it up — is the closest analogue to TAMPURA’s *Find Die* task. Read carefully, the comparison cuts two ways. On time, TAMPURA’s published 57 ± 38 s (the *Partial Observability* column of Table II in [7]) is a *per-step* planning time including simulated controller execution, not a per-episode figure; rather than compare across mismatched units we re-ran the released *Find Die* environment on this thesis’s machine, where it takes 166 ± 85 s per successful episode against our 13.7 s wall-clock (success-only) — roughly an order of magnitude apart, end-to-end, on the same CPU. The published per-step number tells the same story: our matched per-step cost is 5.9 s. On reliability the order runs against us: our holding success rate is 42 %, below the local TAMPURA run’s 55 % (itself consistent with the ≥ 63 % implied by its 0.63 ± 0.30 discounted return, Table I). We are cheaper end-to-end but less reliable, on a marginally easier goal — *Find Die* also requires returning home, which ours does not. Beyond the numbers, the two systems differ in kind.

Both systems sample geometry *online*, inside the planning loop, and both rely on the same determinised, Fast Downward-based classical planner — but they use it for different ends. TAMPURA’s model-learning runs at every control step (Algorithm 2 of [7]). From the current belief it optimistically assumes each action succeeds and runs Fast Downward to propose candidate action sequences that reach the goal; it then samples those actions’ real outcomes to estimate how they branch under uncertainty, building a *probabilistic* sparse MDP $\bar{\mathcal{B}}_{\text{sparse}}$, and finally solves that MDP for an uncertainty- and risk-aware policy. Fast Downward thus chooses *which goal-reaching action sequences are worth considering*; the MDP solver chooses *what to do once their outcomes are uncertain*, weighing each branch’s probability of success. Our planner runs the same kind of determinised search through PDDLStream, but stops there: it reasons over deterministic knowledge literals, executes the chosen plan, and replans when an optimistic assumption proves wrong — there are no transition probabilities to estimate, no

probabilistic MDP, and no policy layer on top. The salient difference is that probabilistic layer — learning the MDP and solving it for a policy — not the symbolic search engine underneath.

This framing changes what the time comparison means. It is not a “semantic-Boxels-beats-TAMPURA” claim. The two figures come from different environments — our tabletop scenes against TAMPURA’s — so the times are not measuring the same problem and benchmark neither system. The planners do also differ in kind — ours runs deterministic classical search and replans, whereas TAMPURA learns a probabilistic sparse MDP and solves it for a policy. What the comparison does show is that a deterministic, replanning-based planner reaches the goal on the closest analogous task without ever estimating the learned sparse-MDP abstraction TAMPURA relies on. Both planners represent the belief with symbolic atoms — including epistemic ones, such as whether an object’s pose is yet known — so the contrast is not whether knowledge is symbolic but how it is used: we reason over it deterministically and recover by replanning, with no transition probabilities to estimate, at the price of no risk-weighting over uncertain outcomes (Section 6.5).

Beyond *how* each system plans, the comparison highlights *what* each system can reason about. Because our planner holds occlusion as explicit symbolic state, line of sight is a planning condition: the sense action requires a clear view of its target region, and the objects that block a view are explicit facts in the planner’s state. Those facts also cover candidate placements: a free cell lying on the camera’s line of sight to a shadow is flagged as view-blocking, so the planner will not set an object down where it would block a region it still needs to observe. TAMPURA’s *Find Die* benchmark—its hidden-object task, the closest analogue to our setting—instead keeps occlusion sub-symbolic: a visibility voxel grid that only weights the success probability of the look action, and a placement sampler that draws object poses with no visibility check, so an object can be placed view-blind. Making space plannable in this way is the point of the representation.

6.4 Threats to Validity

The findings above are conditional on the scenes, the budget, and the oracle perception. Five threats merit explicit mention.

Scene family. All measurements are on tabletop scenes with a fixed overhead camera, a Franka Panda, and 6–9 cm cubes. The adaptive-vs-uniform gap is most pronounced in this regime because free-space cells dominate the uniform partition; in scenes with denser objects or smaller workspaces the ratio could narrow.

Oracle perception. Object detection reads ground-truth poses; a learned detector consuming the rendered RGB-D would surface noise, misclassification, and partial-detection failure modes the current numbers do not reflect.

TAMPURA via published numbers. We compare against TAMPURA’s Table II rather than re-running their code. Algorithmic, implementation, and parameter differences are bundled into a single delta. Both planners are single-threaded, and TAMPURA’s CPU has a marginally higher base clock (2.5 vs 2.0 GHz), so the hardware difference, if anything, slightly favours TAMPURA rather than us. Re-running both systems on a shared scene file would strengthen the comparison.

Resolution regime. The free-space leaf size behaves as a flat knob across the full $\sim 4\times$ range tested. Below *auto_cell*, forcing a 5 cm floor changes nothing (the mbs0.05 null finding); above it, coarsening the free-space leaf to $1.5\times$ and $2\times$ *auto_cell* leaves the success rate essentially unchanged — identical seed-for-seed on stack, within seed noise on holding, and only a mild ~ 4 pp drop on find-and-tray-stack — while shrinking the free-space partition by 30–40 %. The cause is structural rather than numerical: the merge step combines free-space cells whenever their merged region is convex, so the partition reflects the convex structure of the free space rather than the leaf size itself. *auto_cell* (the largest visible footprint plus a 1 cm margin) is therefore a convenient choice rather than a tuned constant, and semantic’s compactness advantage over uniform is a structural property of the adaptive partition, not an artefact of resolution tuning. Neither result generalises to scenes where placements are constrained by sub-5 cm geometry; constructing such a scene is the natural follow-up (Section 7.1).

Bounded give-up. The three-strike give-up on still-blocked shadows is unsound for shadows that are physically unreachable. The run report distinguishes shadows observed empty from shadows given up this way, and the evaluation counts the give-up case as a failure; Section 6.5 expands on the consequence.

The Limitations and Accepted Simplifications below catalogue the design decisions that fix the scope of these results.

6.5 Limitations and Accepted Simplifications

The system evaluated in this thesis embodies a set of deliberate simplifications. Each was adopted to keep the implementation tractable and to isolate the contribution — semantic Boxel discretization and its integration with partially observable deterministic TAMP — from concerns that are orthogonal to it. This subsection consolidates these simplifications so that the empirical claims of Chapter 5 are read with the correct scope.

Perception. The scene contains a single fixed overhead camera that observes the whole table. Its rendered RGB-D images are not processed, however: object detection is performed by an oracle that queries ground-truth object poses directly from the simulator and uses ray-casts to determine which objects are visible from the camera’s viewpoint. There is therefore no learned detector and no sensor-noise model. This isolates the planning contribution from perception error — a real perception module consuming the camera images could replace the oracle without altering the planning architecture. Because the camera is fixed, the robot never moves to a sensing configuration; this is sufficient for the tabletop scenes studied here, where every shadow is visible from one viewpoint, but a robot-mounted sensor would be required for occlusions not visible from any fixed vantage point.

Manipulation and execution. Grasping is constraint-based: a held object is rigidly welded to the gripper, so objects cannot slip or be dropped in transit. Only top-down grasps are sampled, at a single fixed height offset. After a place or stack, execution adds a small hardcoded vertical lift before returning; this lift is invisible to the planner — it produces no PDDL state change and moves only along the gravity axis — and exists solely to give the sampling-based motion planner a collision-free start configuration for the next move. Collision events detected during execution are logged but do not trigger an automatic replan.

Spatial representation. Free-space cells are merged only across exactly aligned shared faces; a full semantic merge — matching view-blocking, observability, and ordering — was not implemented, so the partition contains more Boxels than strictly necessary. When an occluder is relocated, its previously computed shadow geometry and the shadow-blocker map are not recomputed at its new position. This rests on the assumption that the world is static apart from the objects the robot itself moves: nothing else shifts, so the only outdated geometry belongs to the relocated object, whose new pose the planner tracks explicitly and replans from.

Planning. The primary goal studied is holding; `stack` and `find-and-tray-stack` are shipped extensions that broaden goal diversity. A planner-search bias concerns the cost model: the domain gives the `stack` action a cost of 2 while every other action costs 1. This is a hand-tuned bias rather than a claim about stacking—because the planner minimises total plan cost and cannot weigh one action’s outcome likelihood against another’s, the higher cost is what keeps it from substituting a `stack` as a cheap rescue when a `place` onto a free Boxel fails its motion check, forcing it to exhaust placement options first. It is a deterministic stand-in for the probabilistic action selection the POD layer does not perform; `stack` remains available whenever the goal requires it. More fundamentally, when a shadow returns *still-blocked* three times in succession the execution loop marks it as not containing the target and proceeds. The shadow is never directly observed, so for a physically unreachable shadow this violates the invariant that belief should reflect observation. It is accepted to keep the sense-plan-act

loop bounded; the run report distinguishes shadows observed empty from shadows given up this way, and the evaluation counts the give-up case as a failure. The principled remedy — re-grounding the planner’s view-blocking atoms from current physics — is left to future work (Section 7.1).

Parameters. Several geometric constants — grasp height offsets, approach and lift heights, the octree’s minimum leaf size, the camera pose — are tuned for the specific table, robot, and object set used here and would not transfer unchanged to a different setup. Generalizing them, for example by deriving grasp parameters from object geometry, is orthogonal to the core contribution and is noted as future work.

Simplifications disclosed in the approach. Two further simplifications are stated where they are introduced rather than here. The conceptual belief model uses two Know-If literals, $K(p)$ and $K(\neg p)$; the implemented domain collapses these into a single Know-If fluent paired with a value-carrying fluent, an equivalent encoding for the scenarios studied (Section 4.3). The sense action is modeled optimistically — it assumes the target is found — with reactive replanning when execution reveals otherwise, because PDDLStream and FastDownward do not support contingent planning (Chapter 4).

7 Conclusion

This thesis presented a novel approach to Task and Motion Planning under partial observability by integrating a semantic spatial abstraction, termed Boxels, with a Partially Observable Deterministic planning framework. At the core of the methodology is adaptive semantic discretization: it partitions the workspace into Boxels, concentrating detail on the objects and occluded regions that matter for the task. By modeling belief states over these Boxels, our system lets a PDDLStream-based TAMP framework plan around uncertain object locations and occlusions.

The Semantic POD-TAMP model formalizes this integration, enabling the planner to generate and execute plans that interleave information-gathering and manipulation actions. The use of K-literals to represent knowledge within the PDDL state allows for a direct and conceptually simple integration with existing POD planners. As demonstrated, this enables the robot to reason about what it knows and where uncertainty lies, and to take actions to resolve that uncertainty in a targeted manner.

This work contributes a representation in which occlusion is first-class symbolic state, letting a TAMP framework handle partial observability without enumerating an exhaustive state space and without the sub-symbolic visibility grid of a POMDP-based TAMP baseline such as TAMPURA’s *Find Die* benchmark. The evaluation isolates this representation by ablation—the same planner on the adaptive partition against a uniform grid at the same cell scale—and shows on three tabletop goals that it produces an order of magnitude fewer cells, lowering per-call planning time accordingly and solving find-and-tray-stack scenes that the uniform baseline effectively cannot (39.8 % vs 1.3 % success). The semantic+mbs0.05 variant—forcing a 5 cm floor on the free-space leaf size—is indistinguishable from the default in the regime tested, indicating that placement is bottlenecked on object footprint rather than free-space resolution at these scene scales. Against the TAMPURA POMDP-based framework, on the holding task — the closest analogue to TAMPURA’s *Find Die* — our planner completes an episode in a mean 13.7 s wall-clock; re-running TAMPURA’s released *Find Die* on the same hardware takes 166 s per successful episode (its published Table II time of 57 s is a per-step figure, not per-episode), so we are roughly an order of magnitude cheaper end-to-end, though we succeed less often (42 % vs 55 % locally); the comparison is architectural rather than a hardware benchmark, since the difference reflects the planning machinery: TAMPURA learns and solves a probabilistic sparse MDP each step, while our deterministic knowledge-literal planner relies on replanning. The remaining directions for extending this work are outlined below.

7.1 Future Work

The simplifications consolidated in Section 6.5 each suggest a concrete next step. The following directions would extend the system beyond the tabletop scenes evaluated here without altering its core architecture.

Learned perception. The most direct extension is to replace the oracle detector with a learned perception module that operates on the rendered RGB-D images rather than ground-truth simulator poses. Because the oracle is confined behind a narrow detection interface, such a module could be substituted without changing the Boxel discretization or the planning layer, turning the architecture’s perception-agnosticism from a claim into a demonstrated property.

Active sensing with a robot-mounted camera. The current fixed overhead camera observes the whole table at once. A robot-mounted sensor would require the planner to reason about *where to look*: a sensing-configuration stream that computes a robot pose from which a target region becomes visible, integrated as an additional precondition on the sense action. This is the prerequisite for scenes with occlusions not visible from any single fixed viewpoint.

Discovering objects beyond anticipated shadows. When the robot senses a shadow region, any object hidden there is discovered and added to the representation. The scenes studied here are arranged so that every place an object can hide is one such anticipated shadow region. Extending the approach to objects hidden inside concave geometry — such as a container or cupboard, whose interior is itself a hidden region that may hold further objects and occlusions — would require the partitioning to recurse: re-running over the newly revealed interior and recursively bounding whatever it exposes.

Full semantic free-space merge. Free-space cells are presently merged only across exactly aligned shared faces. A full semantic merge that also matches view-blocking, observability, and back-to-front ordering would yield a coarser, more task-relevant partition with fewer Boxels, reducing the planner’s initial-state size.

A PDDL Specification

This appendix lists the complete PDDL domain and stream definitions used by the PDDLStream planner throughout the thesis. The domain (Section A.1) declares the predicates, derived predicates, and the five actions *sense*, *move*, *pick*, *place*, and *stack*. The streams (Section A.2) declare the four samplers that PDDLStream invokes lazily during search: *sample-grasp*, *plan-motion*, *compute-kin*, and *compute-stack-kin*.

A.1 Domain

```
1  ;; =====
2  ;; Semantic Boxel TAMP Domain - PDDLStream Compatible (Untyped)
3  ;; =====
4  ;;
5  ;; Models the robot's CAPABILITIES (sense, move, pick, place, stack) rather
6  ;; than any specific scenario. Scenario-specific spatial relationships
7  ;; (e.g., which objects block which regions) are expressed as problem-level
8  ;; init facts using generic predicates like blocks_view_at.
9  ;;
10 ;; Uses derived predicates for visibility: blocks_view and view_clear are
11 ;; computed automatically from object positions – actions only need to
12 ;; update spatial state (obj_at_boxel), not view predicates.
13 ;;
14 ;; Uses Know-If fluents for partial observability.
15 ;; Object types are encoded as predicates: (Boxel ?x), (Obj ?x), etc.
16
17 (define (domain boxel-tamp)
18   (:requirements :strips :equality :derived-predicates :conditional-effects
19     :action-costs)
20
21   (:predicates
22     ;; --- Type predicates ---
23     (Boxel ?x)
24     (Obj ?x)
25     (Config ?x)
26     (Trajectory ?x)
27     (Grasp ?x)
28
29     ;; --- Boxel classification ---
30     (is_shadow ?b) ; Region not directly visible from the camera
31     (is_object ?b) ; Physical object (can be picked, placed, stacked)
32     (is_free_space ?b) ; Known empty space
```

```

32
33 ;; --- Visibility geometry (static) ---
34 (blocks_view_at ?obj ?b ?region) ; When ?obj is at ?b, it blocks view to ?region
35
36 ;; --- Visibility state (derived – DO NOT use in action effects) ---
37 (blocks_view ?obj ?region) ; ?obj currently blocks view to ?region
38 (view_blocked ?region) ; Some object blocks view to ?region
39 (view_clear ?region) ; No object blocks view to ?region
40
41 ;; --- Ground truth (actual world state) ---
42 (obj_at_boxel ?o ?b) ; Object ?o is physically at boxel ?b
43
44 ;; --- Know-If fluent (do we know the value?) ---
45 (obj_at_boxel_KIF ?o ?b) ; We know whether ?o is at ?b (true or false)
46
47 ;; --- Robot state ---
48 (at_config ?q)
49 (handempty)
50 (holding ?o)
51 (obj_pose_known ?o)
52
53 ;; --- Stream certified facts ---
54 (valid_grasp ?o ?g) ; Grasp ?g valid for object ?o
55 (motion ?q1 ?q2 ?t) ; Trajectory ?t from ?q1 to ?q2
56 (kin_solution ?o ?b ?g ?q) ; Config ?q for picking ?o from ?b with ?g
57 (config_for_boxel ?q ?b) ; Config ?q targets boxel ?b (EE inside ?b)
58 (boxel_fits ?o ?b) ; Boxel ?b is large enough to contain ?o (place
    dest or sense region)
59 (on_surface ?b) ; Boxel ?b rests on a support surface (table)
60
61 ;; --- Stacking ---
62 ;; (on ?o ?support) means ?o sits directly on top of ?support.
63 ;; (on_table ?o) is the explicit "?o rests on the table" counterpart.
64 ;; (clear ?o) means nothing is stacked on ?o.
65 (on ?o1 ?o2)
66 (on_table ?o) ; ?o sits directly on the table
67 (clear ?o)
68 (is_tray ?o) ; ?o is the fixed-base tray support
69 (stack_kin ?o ?on_obj ?g ?q) ; IK config ?q to place ?o on top of ?on_obj
70 )
71
72 ;; =====
73 ;; ACTION COSTS: bias planner toward place over stack
74 ;; =====
75 ;; All actions cost 1 except stack (cost 2). Without this, the planner
76 ;; treated stack and place as equally cheap and chose stack as a "rescue"
77 ;; whenever motion planning to a particular free boxel failed, instead of
78 ;; trying other free boxels. Higher stack cost forces the search to

```

```

79  ;; exhaust place destinations before falling back to stack.
80  (:functions (total-cost))
81
82  ;; =====
83  ;; DERIVED PREDICATES: Visibility from object positions
84  ;; =====
85  ;; blocks_view_at is a static geometric fact in the init state.
86  ;; blocks_view is derived: true when an object is currently at a position
87  ;; that geometrically blocks a view corridor.
88  ;; view_clear is derived via stratified negation: true when no object
89  ;; blocks the view.
90
91  (:derived (blocks_view ?obj ?region)
92    (exists (?b)
93      (and (obj_at_boxel ?obj ?b)
94        (blocks_view_at ?obj ?b ?region))))
95
96  (:derived (view_blocked ?region)
97    (exists (?obj)
98      (blocks_view ?obj ?region)))
99
100  (:derived (view_clear ?region)
101    (and (Boxel ?region)
102      (is_shadow ?region)
103      (not (view_blocked ?region))))
104
105  ;; =====
106  ;; SENSE: Observe a region to check for an object
107  ;; =====
108  ;; Requires clear line of sight to the region.
109  ;; Uses the fixed scene camera – no robot positioning needed.
110  ;;
111  ;; Optimistic effect: assume the target is found. If execution reveals
112  ;; otherwise, the outer loop marks the shadow empty and replans, which
113  ;; avoids contingent planning that PDDLStream and FastDownward do not
114  ;; support.
115  (:action sense
116    :parameters (?o ?region)
117    :precondition (and
118      (Obj ?o)
119      (Boxel ?region)
120      (view_clear ?region)
121      (boxel_fits ?o ?region) ; ?o physically fits inside ?region
122      (not (obj_at_boxel_KIF ?o ?region)) ; Only sense if unknown
123    )
124    :effect (and
125      (obj_at_boxel_KIF ?o ?region) ; Now we know
126      (obj_at_boxel ?o ?region) ; OPTIMISTIC: assume found

```

```

127     (obj_pose_known ?o)
128     (increase (total-cost) 1)
129   )
130 )
131
132 ;; =====
133 ;; MOVE: Move robot from one configuration to another
134 ;; =====
135 ;; ?b is the destination boxel – the stream that produced ?q2 certifies
136 ;; that the end-effector at ?q2 is within boxel ?b.
137 (:action move
138   :parameters (?q1 ?q2 ?b ?t)
139   :precondition (and
140     (Config ?q1)
141     (Config ?q2)
142     (Boxel ?b)
143     (Trajectory ?t)
144     (at_config ?q1)
145     (config_for_boxel ?q2 ?b)
146     (motion ?q1 ?q2 ?t)
147   )
148   :effect (and
149     (at_config ?q2)
150     (not (at_config ?q1))
151     (increase (total-cost) 1)
152   )
153 )
154
155 ;; =====
156 ;; PICK: Pick up an object from a boxel
157 ;; =====
158 ;; Must KNOW object is there (KIF=true AND at=true).
159 (:action pick
160   :parameters (?o ?b ?g ?q)
161   :precondition (and
162     (Obj ?o)
163     (Boxel ?b)
164     (Grasp ?g)
165     (Config ?q)
166     (handempty)
167     (at_config ?q)
168     (clear ?o) ; can't pick an object with something stacked
169                 on top
170     (obj_at_boxel_KIF ?o ?b) ; Must know
171     (obj_at_boxel ?o ?b) ; Must be there
172     (kin_solution ?o ?b ?g ?q)
173   )
174   :effect (and

```

```

174     (holding ?o)
175     (not (handempty))
176     (not (obj_at_boxel ?o ?b))
177     (not (on_table ?o))          ; picking lifts ?o off the table
178                                 ; (idempotent if ?o was on a stack instead)
179     ;; If ?o was stacked on ?x, picking ?o makes ?x clear again. For
180     ;; holding-goal runs (no (on ...) facts) this grounds to a no-op.
181     (forall (?x)
182       (when (on ?o ?x)
183         (and (not (on ?o ?x)) (clear ?x))))
184     (increase (total-cost) 1)
185   )
186 )
187
188 ;; =====
189 ;; PLACE: Place an object in a boxel
190 ;; =====
191 ;; Destination must be free space.
192 (:action place
193   :parameters (?o ?b ?g ?q)
194   :precondition (and
195     (Obj ?o)
196     (Boxel ?b)
197     (Grasp ?g)
198     (Config ?q)
199     (holding ?o)
200     (at_config ?q)
201     (is_free_space ?b)
202     (on_surface ?b)
203     (boxel_fits ?o ?b)
204     (kin_solution ?o ?b ?g ?q)
205   )
206   :effect (and
207     (handempty)
208     (obj_at_boxel ?o ?b)
209     (obj_at_boxel_KIF ?o ?b)
210     (on_table ?o)          ; place lands ?o on the table
211     (not (holding ?o))
212     (not (is_free_space ?b))
213     (increase (total-cost) 1)
214   )
215 )
216
217 ;; =====
218 ;; STACK: Place the held object on top of another object
219 ;; =====
220 ;; Mirrors place but the destination is an OBJECT boxel rather than a
221 ;; FREE_SPACE boxel. ``stack_kin`` is certified by compute-stack-kin,

```

```

222  ;; which derives the EE target from ?on_obj's CURRENT top z + the held
223  ;; object's half-height; this is computed each time the planner is
224  ;; invoked so multi-step stacks tolerate per-step settling.
225  ;;
226  ;; (clear ?o) becomes true on the newly-stacked object so it is itself
227  ;; available as a support for a subsequent stack action.
228  (:action stack
229    :parameters (?o ?on_obj ?g ?q)
230    :precondition (and
231      (Obj ?o)
232      (Obj ?on_obj)
233      (Grasp ?g)
234      (Config ?q)
235      (holding ?o)
236      (at_config ?q)
237      (clear ?on_obj)
238      (stack_kin ?o ?on_obj ?g ?q)
239    )
240    :effect (and
241      (handempty)
242      (on ?o ?on_obj)
243      (clear ?o)
244      (obj_at_boxel ?o ?o) ; OBJECT boxel ID equals object name (a stacked
245        object's boxel is itself)
246      (obj_at_boxel_KIF ?o ?o)
247      (not (holding ?o))
248      (not (clear ?on_obj))
249      (increase (total-cost) 2)
250    )
251  )

```

Listing A.1: The Semantic Boxel TAMP domain.

A.2 Streams

```

1  (define (stream boxel-streams)
2
3  ;; Sample grasp poses for an object
4  (:stream sample-grasp
5    :inputs (?o)
6    :domain (Obj ?o)
7    :outputs (?g)
8    :certified (and (Grasp ?g) (valid_grasp ?o ?g))
9  )
10
11  ;; Plan motion between configurations
12  (:stream plan-motion

```

```

13   :inputs (?q1 ?q2)
14   :domain (and (Config ?q1) (Config ?q2))
15   :outputs (?t)
16   :certified (and (Trajectory ?t) (motion ?q1 ?q2 ?t))
17 )
18
19 ;; boxel_fits is NOT a stream --- the facts are emitted in the initial
20 ;; state. As a test stream it forced adaptive search to re-evaluate
21 ;; every (object, free_boxel) pair across skeletons.
22
23 ;; Compute IK solution for picking
24 (:stream compute-kin
25   :inputs (?o ?b ?g)
26   :domain (and (Obj ?o) (Boxel ?b) (valid_grasp ?o ?g))
27   :outputs (?q)
28   :certified (and (Config ?q) (kin_solution ?o ?b ?g ?q) (config_for_boxel ?q ?b))
29 )
30
31 ;; Compute IK for stacking ?o on top of ?on_obj.
32 ;; Reads ?on_obj's CURRENT AABB top + held object half-height to derive
33 ;; the EE target. Certifies config_for_boxel against ?on_obj so the
34 ;; preceding move action delivers the arm to the support's OBJECT boxel.
35 (:stream compute-stack-kin
36   :inputs (?o ?on_obj ?g)
37   :domain (and (Obj ?o) (Obj ?on_obj) (valid_grasp ?o ?g))
38   :outputs (?q)
39   :certified (and (Config ?q) (stack_kin ?o ?on_obj ?g ?q) (config_for_boxel ?q
40     ?on_obj))
41 )

```

Listing A.2: Stream definitions.

List of Acronyms

AI	Artificial Intelligence
CR	Critical Region
FOND	Fully-Observable Non-Deterministic
IK	Inverse Kinematics
KIF	Know-If Fluent
LAO*	Loop AO* Search
LLM	Large Language Model
MDP	Markov Decision Process
PDDL	Planning Domain Definition Language
POD	Partially Observable Deterministic
POMDP	Partially Observable Markov Decision Process
RBVD	Region-Based Voronoi Diagram
STRIPS	Stanford Research Institute Problem Solver
TAMP	Task and Motion Planning
TAMPURA	TAMP with Uncertainty and Risk Awareness
VLM	Vision-Language Model

List of Symbols

State-Space Planning Model

M	deterministic state-space model $\langle S, s_0, S_G, Act, A, f \rangle$
S	set of states
s	a state
s_0	initial state
s'	successor state of s
S_G	set of goal states
Act	set of all (ground) actions
$A(s)$	actions applicable in state s
a	an action
$f(a, s)$	deterministic state-transition function
σ	a plan, i.e. a sequence of actions
P	a PDDL problem

Belief and Knowledge

b	belief state (set of possible states)
p	a proposition (ground atom)
L	a literal
$K(\cdot)$	knowledge (K-literal) operator: its argument is known to hold
$\neg$	logical negation

Semantic POD-TAMP Model

$\mathcal{M}$	semantic POD-TAMP model $\langle \mathcal{S}, \mathcal{S}_0, \mathcal{S}_G, \mathcal{A}, f, \mathcal{O} \rangle$
$\mathcal{S}$	set of abstract (belief) states
$\mathcal{S}_0$	set of initial abstract states
$\mathcal{S}_G$	set of goal states
$\mathcal{A}$	set of abstract action schemas
$\mathcal{O}$	sensor model
$\mathcal{AP}$	set of atomic propositions
b_0	initial belief
b_g	goal belief

Semantic Boxels

$\mathcal{BX}_{\text{set}}$	set of Semantic Boxels
obj	a manipulable object

$Boxel$	a Boxel (workspace region)
$InBoxel$	predicate: object obj is in $Boxel$
N	number of candidate regions to search

TAMPURA Comparison

$\bar{\mathcal{B}}_{\text{sparse}}$	probabilistic sparse MDP learned by TAMPURA
-------------------------------------	---

List of Figures

1.1	A 7-DOF Franka Panda at a cluttered tabletop. The target—the cyan cube behind the green cube—is hidden from the overhead camera and must be located before it can be retrieved. The corner panels are the fixed overhead camera’s synthetic RGB (top) and depth (bottom) views.	2
2.1	Example of an octree storing free (shaded white) and occupied (black) cells. The volumetric model is shown on the left and the corresponding tree representation on the right.	9
4.1	Adaptive Semantic Discretization. (a) Initial scene: the camera viewpoint and the detected objects on the table, with a fully occluded target (cyan) sitting in the shadow behind the largest occluder. (b) Each object is bounded by a dedicated object Boxel. (c) The volume each object hides from the camera becomes an occlusion (shadow) Boxel, swept from the object’s back face out to the workspace boundary along the lines of sight (dashed) and rising to the object’s height. (d)–(f) The remaining free space is partitioned by an octree: from a single cell spanning the whole workspace (d), only cells containing both free and occupied space are recursively subdivided—in all three dimensions, so the open space above the objects stays a few large cells (e)—and adjacent free cells are then greedily merged across aligned faces into larger convex Free Space Boxels (f).	17
4.2	The same partition applied to a PyBullet scene. The cyan wireframe cells are the Free Space Boxels filling the workspace; at the cube cluster, object bounding cuboids enclose the visible cubes and red wireframe Boxels mark the volumes they occlude.	17
4.3	Adaptive semantic partition (left) versus uniform-grid baseline (right) on a matching scene. The adaptive partition produces a small set of task-relevant cells; the uniform grid covers the whole workspace at the same cell size, producing an order of magnitude more cells.	19
4.4	The sense action targeting a shadow region in PyBullet. The arm has been moved to a sensing configuration over a labelled shadow; the in-simulator overlay labels each object and its shadow region. This is the configuration in which the sense action’s (view_clear ?region) precondition is checked and its optimistic (obj_at_boxel ?o ?region) effect is asserted.	23

5.1	The three evaluated goals. holding: retrieve a single target hidden behind occluders. stack: build a tower of cubes to a target height. find-and-tray-stack: locate hidden cubes and stack them on a tray.	24
5.2	The occluder-count difficulty axis n_{occ} , shown at low, medium, and high settings. More occluders crowd the target and occlude a larger share of the workspace, so the planner must reason over more shadow Boxels.	25
5.3	An evaluation scene observed by the fixed overhead camera, shown here from an oblique simulator viewpoint; the corner panels are the camera’s synthetic RGB (top) and depth (bottom) views. Object detection is an oracle that reads ground-truth object poses from the simulator rather than processing the rendered RGB-D images. A future learned detector would consume those images in place of the oracle, without changing the planning architecture.	26
5.4	Cumulative solve rate vs wall-clock budget (log axis) per (goal, variant). Almost all solves happen within ~ 200 s; uniform starts later and plateaus lower on every goal.	29
5.5	Success rate vs n_{occ} on the holding goal. Adaptive variants (semantic, semantic+mbs0.05) are within seed noise of each other; uniform sits ~ 10 pp lower. Each bar is the mean success rate over the ~ 80 – 100 seeds at that difficulty (here and in Figures 5.6 and 5.7).	29
5.6	Success rate vs stack height h on the stack goal. Adaptive variants degrade gracefully; uniform collapses past $h = 2$	30
5.7	Success rate vs n_{occ} on the find-and-tray-stack goal. Uniform fails almost entirely; semantic+mbs0.05 visibly degrades with n_{occ} , contrary to its intent. . .	30
5.8	Mean success-only planning time vs n_{occ} on the holding goal. Adaptive variants scale gracefully; uniform’s cost grows much faster.	31
5.9	Mean Boxel count by category, holding goal. Adaptive variants produce an order of magnitude fewer cells than the uniform baseline.	32
5.10	Mean Boxel count vs free-space leaf size for the find-and-tray-stack and holding goals (semantic variant; total and free-space cells). Each bar group is one resolution arm, labelled by its mean effective leaf size. The count is flat at and below the auto-set leaf size and falls 30–40 % as the leaf is coarsened above it. .	32
5.11	Per-episode time on holding (same hardware): our wall-clock (semantic variant; mean $\pm$ std, success-only; right-skewed, median 8.5 s) and real TAMPURA <i>Find Die</i> re-run locally (mean $\pm$ std over successful episodes; 20 seeds, 55 % success). Both are per-episode times including simulated execution. TAMPURA’s published 57 s (Table II) is a per-step figure and is not plotted; see the text for the per-step reconciliation. On reliability the order reverses, holding-semantic succeeding on 42 % of episodes against TAMPURA’s local 55 %.	34

- 5.12 Stacked exit-reason breakdown per (goal, variant). `planner_failed` here also absorbs the `physics_mismatch`, `drop_failed` and `all_searched` exits (all cases where the planned solution ultimately failed); it is the dominant failure mode on holding and find-and-tray-stack, while on stack the semantic variant instead fails mainly via `replan_limit` (the second-largest slice after success). 35

List of Tables

5.1	Headline success rate and mean planning time per (goal, variant). Planning time is success-only.	28
-----	--	----

List of Listings

4.1	PDDL fluents for belief over Boxels	21
4.2	The implemented sense action	21
A.1	The Semantic Boxel TAMP domain.	44
A.2	Stream definitions.	49

List of References

- [1] A. Albore, H. Palacios, and H. Geffner, “A translation-based approach to contingent planning,” in *IJCAI*, vol. 9, 2009, pp. 1623–1628.
- [2] Y. Bai, Z. Wang, Y. Liu, K. Luo, Y. Wen, M. Dai, W. Chen, Z. Chen, L. Liu, G. Li, and L. Lin, “Learning to see and act: Task-aware virtual view exploration for robotic manipulation,” *arXiv preprint arXiv:2508.05186*, 2025.
- [3] W. Bejjani, W. C. Agboh, M. R. Dogar, and M. Leonetti, “Occlusion-aware search for object retrieval in clutter,” in *2021 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, IEEE, 2021, pp. 4678–4685.
- [4] B. Bonet and H. Geffner, “Planning under partial observability by classical replanning: Theory and experiments,” in *Proceedings of the 22nd International Joint Conference on Artificial Intelligence (IJCAI)*, 2011, pp. 1936–1941.
- [5] B. Bonet and H. Geffner, “Flexible and scalable partially observable planning with linear translations,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 28, 2014, pp. 2235–2241.
- [6] E. Coumans and Y. Bai, *PyBullet, a Python module for physics simulation for games, robotics and machine learning*, <http://pybullet.org>, 2021.
- [7] A. Curtis, G. Matheos, N. Gothoskar, V. Mansinghka, J. Tenenbaum, T. Lozano-Pérez, and L. P. Kaelbling, “Partially observable task and motion planning with uncertainty and risk awareness,” *arXiv preprint arXiv:2403.10454v2*, 2024.
- [8] R. E. Fikes and N. J. Nilsson, “Strips: A new approach to the application of theorem proving to problem solving,” *Artificial intelligence*, vol. 2, no. 3-4, pp. 189–208, 1971.
- [9] C. R. Garrett, R. Chitnis, R. Holladay, B. Kim, T. Silver, L. P. Kaelbling, and T. Lozano-Pérez, “Integrated task and motion planning,” *Annual review of control, robotics, and autonomous systems*, vol. 4, no. 1, pp. 265–293, 2021.
- [10] C. R. Garrett, T. Lozano-Pérez, and L. P. Kaelbling, *PDDLStream: Integrating symbolic planners and blackbox samplers via optimistic adaptive planning*, arXiv:1802.08705. Also in Proc. of ICAPS 2020, 2018. arXiv: 1802.08705.
- [11] C. R. Garrett, C. Paxton, T. Lozano-Pérez, L. P. Kaelbling, and D. Fox, “Online replanning in belief space for partially observable task and motion problems,” in *IEEE International Conference on Robotics and Automation (ICRA)*, 2020, pp. 5678–5684.

-
- [12] H. Geffner and B. Bonet, *A concise introduction to models and methods for automated planning*. Morgan & Claypool Publishers, 2013.
 - [13] M. Ghallab, D. Nau, and P. Traverso, *Automated Planning: theory and practice*. Morgan Kaufmann, 2004.
 - [14] N. Gothoskar, M. Ghavami, E. Li, A. Curtis, M. Noseworthy, K. Chung, B. Patton, W. T. Freeman, J. B. Tenenbaum, M. Klukas, and V. K. Mansinghka, *Bayes3D: Fast learning and inference in structured generative models of 3d objects and scenes*, arXiv:2312.08715, 2023. arXiv: 2312.08715.
 - [15] D. Hadfield-Menell, E. Groshev, R. Chitnis, and P. Abbeel, “Modular task and motion planning in belief space,” in *2015 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, IEEE, 2015, pp. 4991–4998.
 - [16] E. A. Hansen and S. Zilberstein, “Lao*: A heuristic search algorithm that finds solutions with loops,” *Artificial Intelligence*, vol. 129, no. 1-2, pp. 35–62, 2001.
 - [17] T. Hofmann and H. Geffner, “Learning generalized policies for fully observable non-deterministic planning domains,” *arXiv preprint arXiv:2404.02499*, 2024.
 - [18] A. Hornung, K. M. Wurm, M. Bennewitz, C. Stachniss, and W. Burgard, “Octomap: An efficient probabilistic 3d mapping framework based on octrees,” *Autonomous robots*, vol. 34, no. 3, pp. 189–206, 2013.
 - [19] L. P. Kaelbling, M. L. Littman, and A. R. Cassandra, “Planning and acting in partially observable stochastic domains,” *Artificial intelligence*, vol. 101, no. 1-2, pp. 99–134, 1998.
 - [20] L. P. Kaelbling and T. Lozano-Pérez, “Integrated task and motion planning in belief space,” *The International Journal of Robotics Research*, vol. 32, no. 9-10, pp. 1194–1227, 2013.
 - [21] Y. Kim, R. Arora, R. Martín-Martín, P. Stone, B. Abbatematteo, and Y. Sung, “Large-language-model-guided state estimation for partially observable task and motion planning,” *arXiv preprint arXiv:2603.03704*, 2026.
 - [22] Y. Ma, Y. Yuan, S. Wu, and H. Yuan, “A task and motion planning framework for partially observable household manipulation scenes,” *Advanced Intelligent Systems*, vol. 7, no. 12, p. 2400897, 2025.
 - [23] D. McDermott, M. Ghallab, A. Howe, C. Knoblock, A. Ram, M. Veloso, D. Weld, and D. Wilkins, “PDDL – the planning domain definition language,” AIPS-98 Planning Competition Committee, Tech. Rep. CVC TR-98-003, 1998.
 - [24] D. Molina, K. Kumar, and S. Srivastava, “Learn and link: Learning critical regions for efficient planning,” in *2020 IEEE International Conference on Robotics and Automation (ICRA)*, 2020, pp. 10 605–10 611.

- [25] C. Muise, V. Belle, and S. A. McIlraith, “Computing contingent plans via fully observable non-deterministic planning,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 28, 2014, pp. 2322–2329.
- [26] T. Pan, R. Shome, and L. E. Kavraki, “Task and motion planning for execution in the real,” *IEEE Transactions on Robotics*, vol. 40, pp. 3356–3371, 2024.
- [27] G. Riegler, A. Osman Ulusoy, and A. Geiger, “Octnet: Learning deep 3d representations at high resolutions,” in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2017, pp. 6620–6629.
- [28] M. S. Saleem, R. Veerapaneni, and M. Likhachev, “A pomdp-based hierarchical planning framework for manipulation under pose uncertainty,” *arXiv preprint arXiv:2409.18775*, 2024.
- [29] N. Shah and S. Srivastava, “Using deep learning to bootstrap abstractions for hierarchical robot planning,” in *Proc. of the 21st International Conference on Autonomous Agents and Multiagent Systems (AAMAS ’22)*, 2022, pp. 1183–1191.
- [30] L. Zhao, W. McClinton, A. Curtis, N. Kumar, T. Silver, L. P. Kaelbling, and L. L. Wong, “Seeing is believing: Belief-space planning with foundation models as uncertainty estimators,” *arXiv preprint arXiv:2504.03245*, 2025.